

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
Entering the Kernel, main

Carmi Merimovich

Tel-Aviv Academic College

September 18, 2017

xv6 kernel programming model

- Like for user mode, there is an **xv6** programming model.
- For user mode, the kernel is the maintainer of the model.
- The kernel is the model maintainer for the kernel programming.
- The kernel is trustable, so the model can include a “don’t do” rules

The xv6 kernel programming model

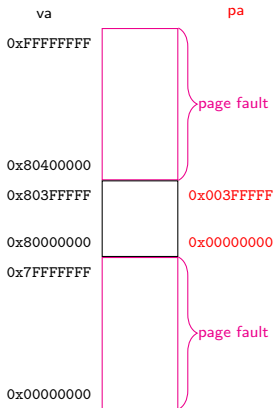
- The kernel is linked into virtual address 0x80100000 and up.
- The code can freely use the general purpose registers: **eax**, **ecx**, **edx**, **ebx**, **esi**, **edi**, **ebp**.
- There is a 4KB stack for the kernel to use. No messing up with it!
- Registers **cs**, **ds**, **ss**, **es** should not be modified.
- (the base and limit are assumed to be zero and max in the above registers).

The REAL entry point on the primary processor is **_start/entry**

entry does the following:

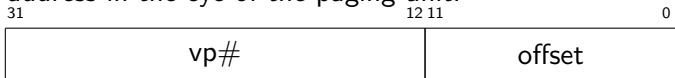
- Builds a temporary page table which maps 4MB of virtual addresses beginning at 0x80000000 to 4MB of physical address beginning at 0x00000000.
- Note the kernel was loaded at physical address 0x0010000 and the programming model is the kernel begins at 0x80100000.
- Starts the paging unit with the above constructed table.
- Points register **esp** to a 4KB stack area.
- (This would be the kernel stack of the primary processor.)
- Jumps to **main**.

Memory after **entry** execution

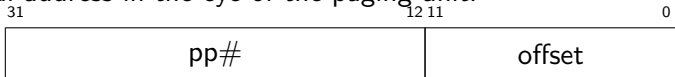


Recall

- Virtual address in the eye of the paging unit:



- Physical address in the eye of the paging unit:



- Page table entry:



Setting up the temporary translation table

The 4MB translation

- 4MB range : 0x00000000-0x003FFFFFFF.
- In ranges of 1MB:

0x800xxxxx \mapsto 0x000xxxxx

0x801xxxxx \mapsto 0x001xxxxx

0x802xxxxx \mapsto 0x002xxxxx

0x803xxxxx \mapsto 0x003xxxxx

- In ranges of 4KB:

0x80000xxx \mapsto 0x00000xxx

0x80001xxx \mapsto 0x00001xxx

0x80002xxx \mapsto 0x00002xxx

0x80003xxx \mapsto 0x00003xxx

\vdots

The rule: $0x800xxxxx \mapsto 0x000xxxxx$

- Translating to pages we have: $0x800xx \mapsto 0x000xx$.
- In table form:

vp#		pp#	
dec	hex	hex	dec
524288	80000	00000	0
524289	80001	00001	1
⋮	⋮	⋮	⋮
524541	800FE	000FE	254
524542	800FF	000FF	255

- Coding it (for an **hypothetical** linear page table):

```
for (i = 0; i < 256; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;
```

The rule: $0x801xxxxx \mapsto 0x001xxxxx$

- The rule for pages is: $0x801xx \mapsto 0x001xx$.
- In table form:

vp#		pp#	
dec	hex	hex	dec
524544	80100	00100	256
524545	80101	00101	257
⋮	⋮	⋮	⋮
524798	801FE	001FE	510
524799	801FF	001FF	511

- Coding it (for linear page table):

```
for (i = 0x100; i < 0x200; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;
```

Hypothetical code for the four rules

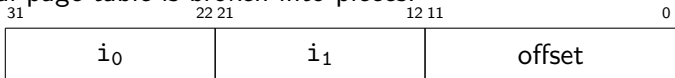
```
for (i = 0x000; i < 0x100; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;  
for (i = 0x100; i < 0x200; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;  
for (i = 0x200; i < 0x300; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;  
for (i = 0x300; i < 0x400; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;
```

Which is equivalent to the following.

```
for (i = 0; i < 1024; i++)  
    pgtbl[0x80000+i] = (i << 12) | 3;
```

Real page table is hierarchical

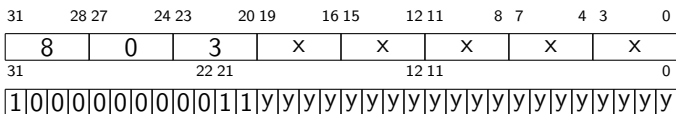
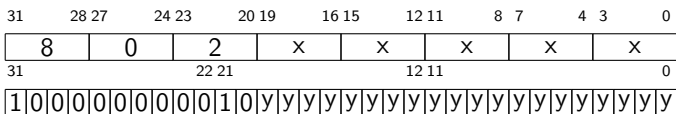
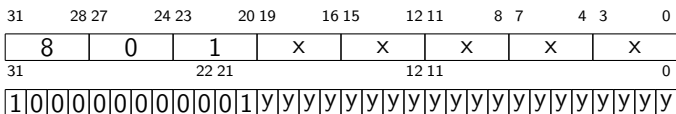
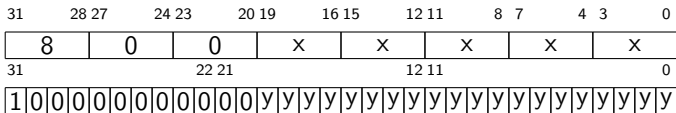
- The real page table is broken into pieces.



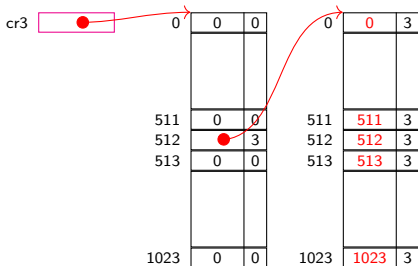

virtual address as viewed by the paging unit

- For each index $0x80000 + i$ we have:
 - Piece number: $(0x80000 + i)/1024$.
 - Index in the piece: $(0x80000 + i)\%1024$.
- In computer lingo, for each index $0x80000 + i$ we have:
 - Piece number: $(0x80000 + i) >> 10$.
 - Index in the piece: $(0x80000 + i) \& 1023$.

Pieces in binary



Pieces of $0x80000 + i$ ($i < 1024$)



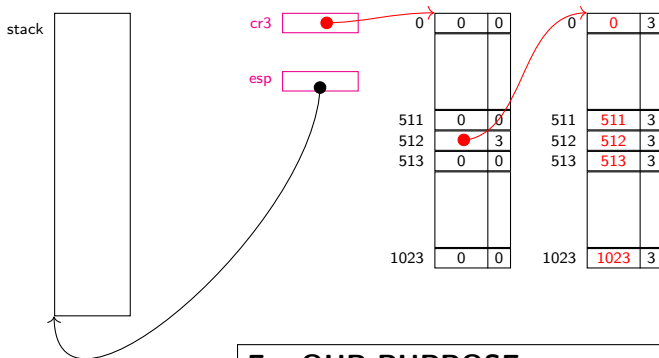
Coding for real table

```
static int
```

```
    pgdir[1024] __attribute__((aligned(4096))),  
    pgtbl[1024] __attribute__((aligned(4096)));
```

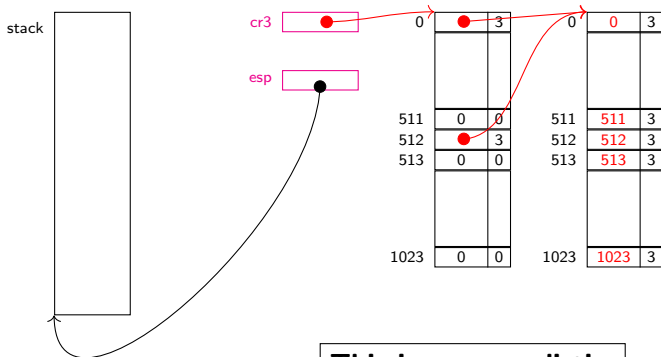
```
memset(pgdir, 4096, 0);  
pgdir[512] = pgtbl | 3;  
for (i = 0; i < 1024; i++) {  
    pgtbl[i] = (i << 12) | 3;  
}  
lcr3(pgdir);  
lcr0(rcr0() | CR0_PG);  
// Paging has gone active  
(Spectacular crash)
```

state when entering main



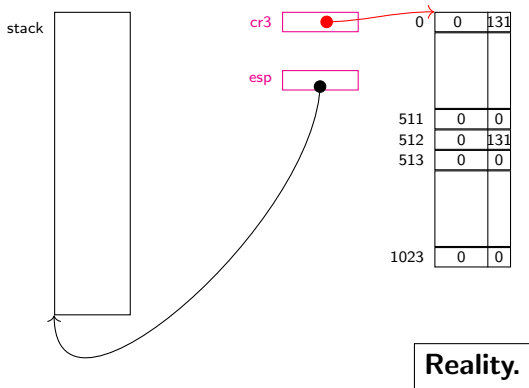
**For OUR PURPOSE,
this is a good enough description,
eventhough it is not the reality.**

state when entering main (elective)



This is more realistic.

state when entering main (elective)



entry.s (elective)

```
.globl _start
_start = V2P_WO(entry)
.globl entry
entry:
    movl %cr4,%eax
    orl $CR4_PSE,%eax
    movl %eax, %cr4
    movl $V2P_WO(entrypgdir),%eax
    movl %eax, %cr3
    movl %cr0, %eax
    orl $CR0_PG|CR0_WP,%eax
    movl %eax,%cr0
```

```
movl $stack+KSTACKSIZE,%esp

mov $main,%eax
jmp *%eax

.comm stack ,KSTACKSIZE
```

main

Since we got here from **entry**, this code is run by the primary processor.

- **main** calls routines initializing the different subsystems.
- We study each initialization together with its subsystem.
- Most subsystems need one initialization per (computer) system.
- Initialization which are needed per-cpu are done in **mpmain**.

The auxiliary processors begins at **entryother**, then proceed to **mpenter**.

main

1216

```
int main(void) {  
    kinit1(end,  
            P2V(4*1024*1024));  
    kvmalloc();  
    mpinit();  
    lapicinit();  
    seginit();  
    picinit();  
    ioapicinit();  
    consoleinit();  
    uartinit();  
    pinit();  
    tvinit();  
    binit();
```

```
    fileinit();  
    iinit();  
    ideinit();  
    if (!ismp) timerinit();  
    startothers();  
    kinit2(P2V(4*1024*1024),  
           P2V(PHYSTOP));  
    userinit();  
    mpmain();  
}
```

The initializations we study now

```
kinit1(end, P2V(4*1024*1024)); // phys page allocation
kvmalloc(); // kernel page table
:
segininit(); // set up segments
:
pinit(); // process table
:
:
kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must copy
userinit(); // first user process
mpmain();
```