

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
Scheduler

Carmi Merimovich

Tel-Aviv Academic College

November 20, 2017

Context

```
1219 kinit1(end, P2V(4*1024*1024)); // phys page allocation
    kvmalloc(); // kernel page table
    :
1222 seginit(); // set up segments
    :
1224 pinit(); // process table
    :
    :
1226 kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must copy
1227 userinit(); // first user process
mpmain();
```

Auxiliary context

- The primary processor begins its C code in **main()**.
- The auxiliary processors begins their C code in **mpenter()**.
- The state on entering either **main()** or **mpenter()** is the same.
- There is a separate stack of each processor.

```
1241 static void mpenter(void) {  
    switchkvm();  
    seginit();  
    lapicinit();  
    mpmain();  
}
```

mycpu()

```
2436 struct cpu* mycpu(void) {  
    int apicid, i;  
  
    if (readeflags() & FL_IF)  
        panic("mycpu_called_with_interrupts_enabled\n");  
  
    apicid = lapicid();  
    for (i = 0; i < ncpu; ++i) {  
        if (cpus[i].apicid == apicid)  
            return &cpus[i];  
    }  
    panic("unknown_apicid\n");  
}
```

mpmain()

```
1252 static void mpmain(void) {  
    cprintf("cpu%d: _starting _%d\n", cpuid(), cpuid());  
    idtinit(); // load idt register  
    xchg(&(mycpu()->started), 1); // tell startothers()  
    scheduler(); // start running processes  
}
```

myproc()

```
2456 struct proc *myproc(void) {  
    struct cpu *c;  
    struct proc *p;  
    pushcli();  
    c = mycpu();  
    p = c->proc;  
    popcli();  
    return p;  
}
```

scheduler

2758

```
void scheduler(void) {  
    struct proc *p;  
    struct cpu *c = mycpu();  
    c->proc = 0;  
    for(;;) { sti();  
        acquire(&ptable.lock);  
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {  
            if (p->state != RUNNABLE) continue;  
  
            c->proc = p;  
            switchvm(p);  
            p->state = RUNNING;  
            swtch(&c->scheduler, p->context);  
            switchkvm();  
  
            c->proc = 0;  
        }  
        release(&ptable.lock);  
    }
```

scheduler() operation

- For each proc struct `p` with state `RUNNABLE` the following is executed:
 - `c->proc = p;`
 - `switchvm()`.
 - `swtch()`.
 - `switchkvm()`.
 - `c->proc=NULL.`

swtch()

3058

```
.globl swtch
swtch:
    movl 4(%esp), %eax
    movl 8(%esp), %edx

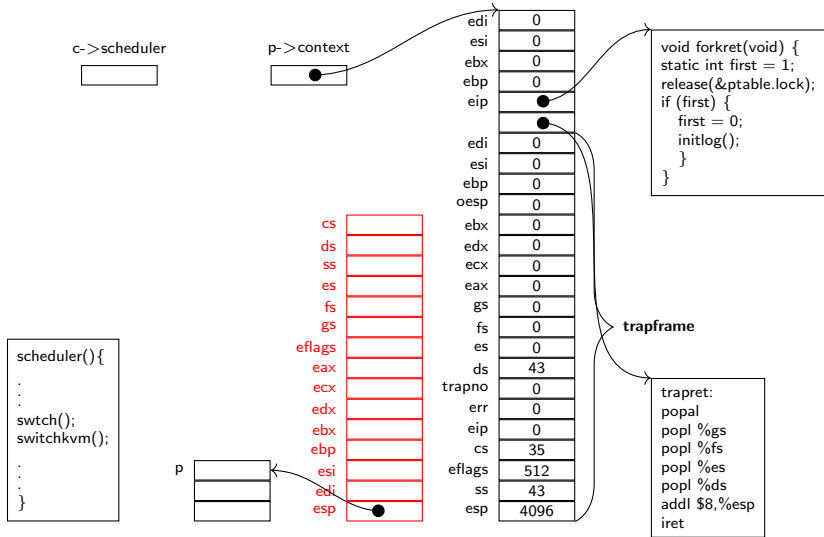
    pushl %ebp
    pushl %ebx
    pushl %esi
    pushl %edi

    movl %esp, (%eax)
    movl %edx, %esp
```

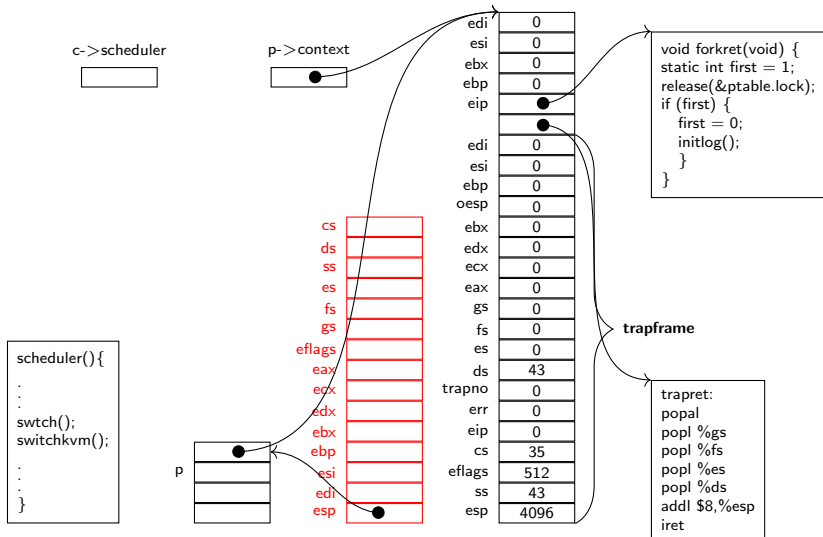
```
    popl %edi
    popl %esi
    popl %ebx
    popl %ebp
    ret
```

The eip field of context is generated by the instruction calling swtch.

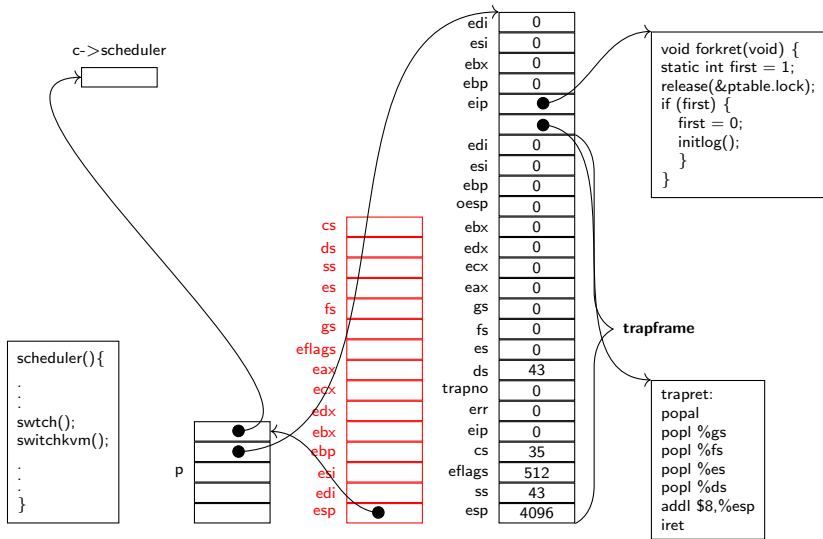
swtch() operation



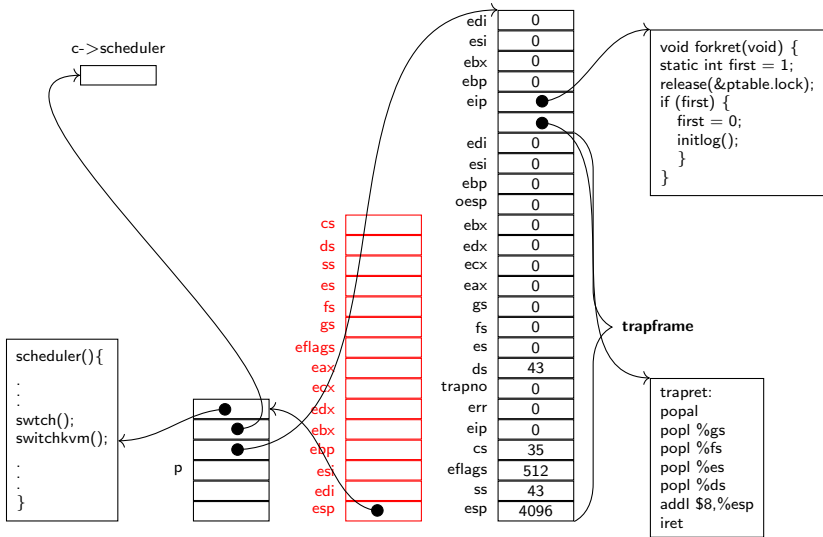
swtch() operation



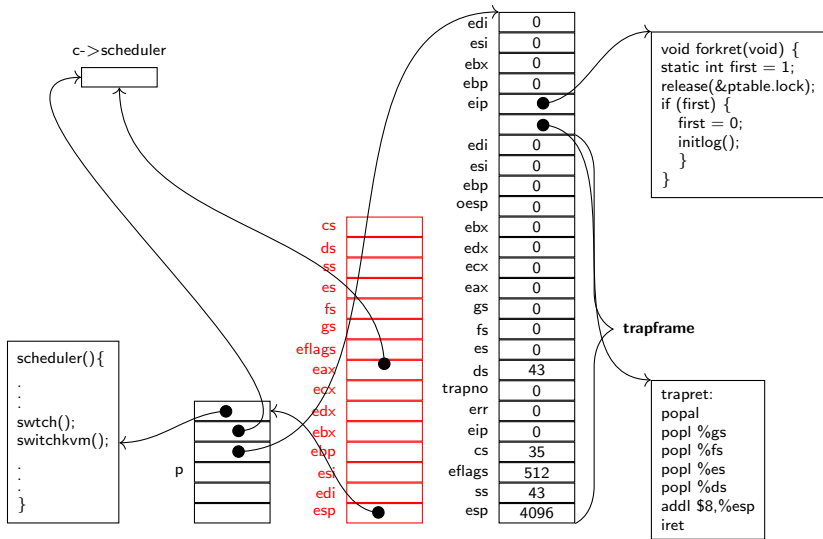
swtch() operation




swtch() operation



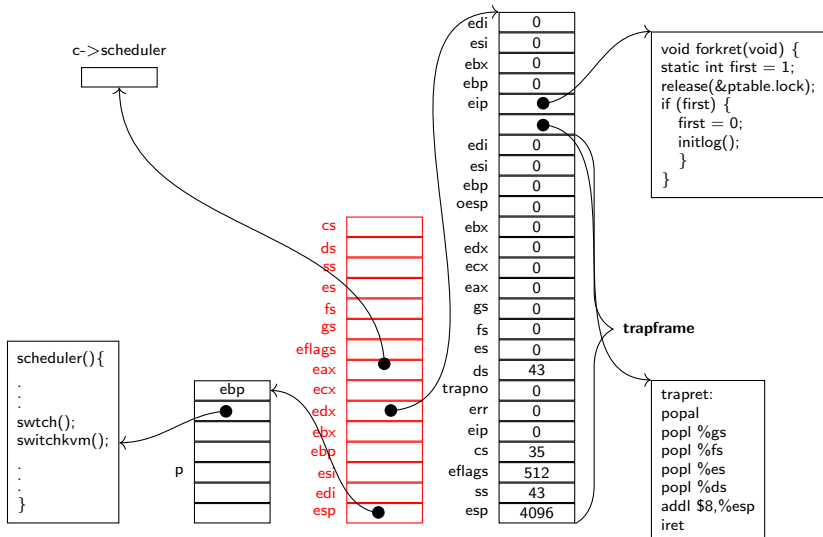
swtch() operation



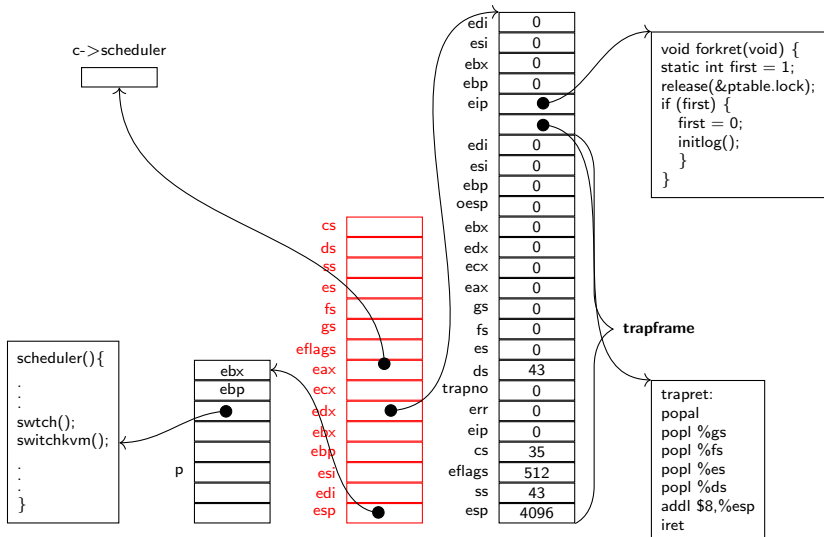
swtch() operation




swtch() operation




swtch() operation



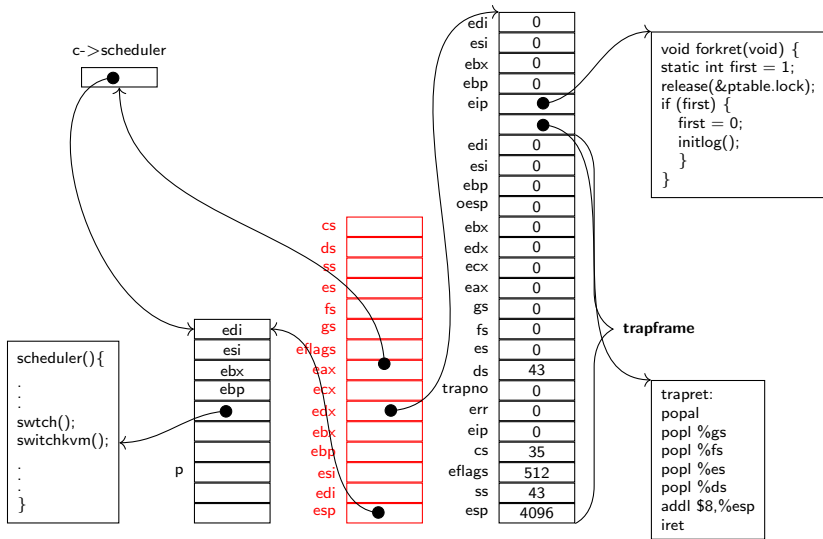
swtch() operation




swtch() operation



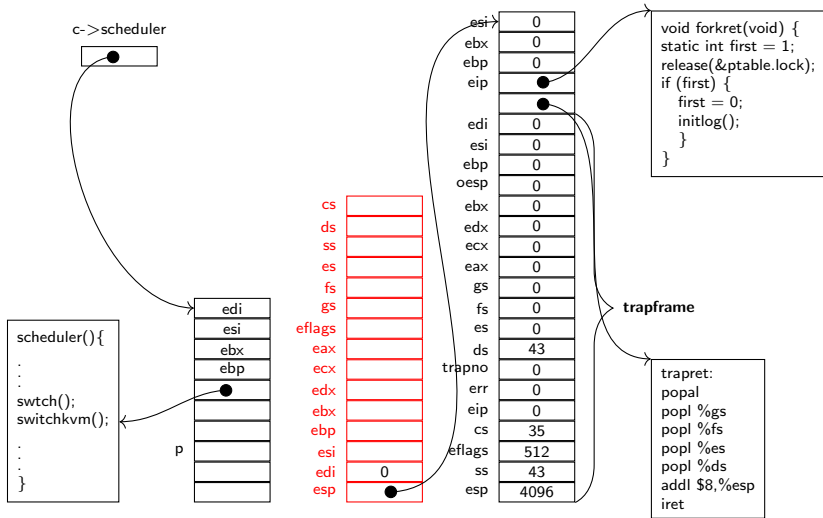
swtch() operation



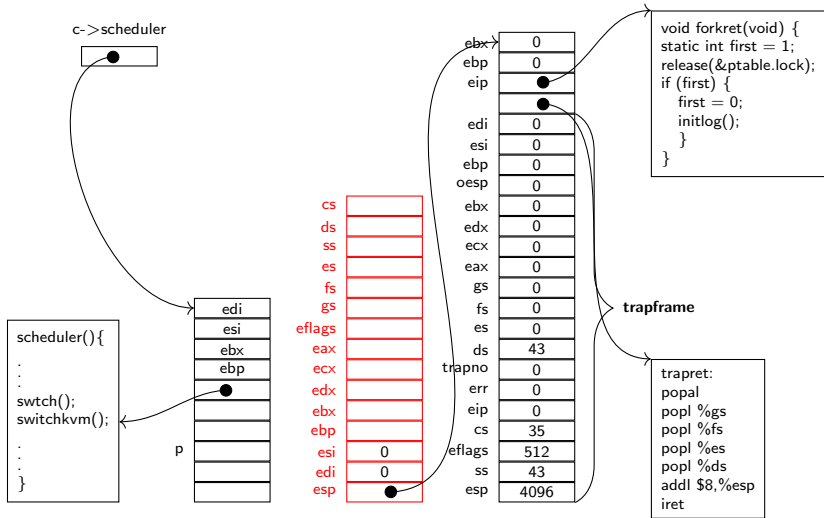
swtch() operation



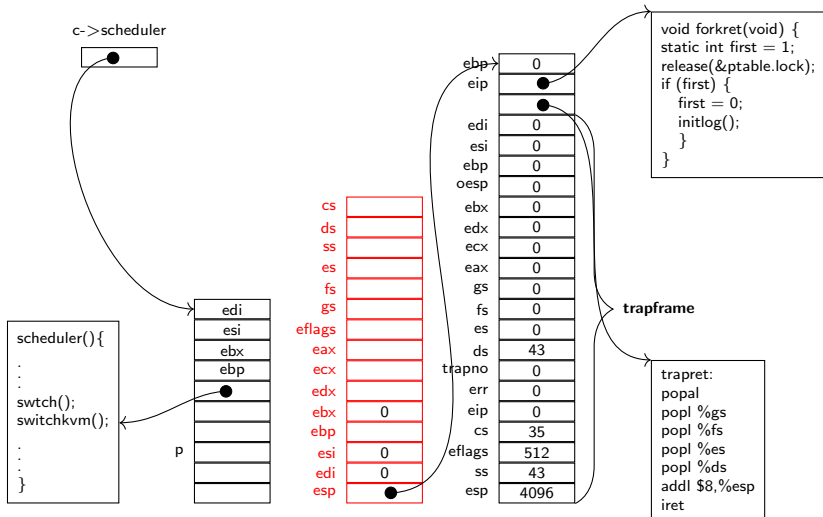
swtch() operation



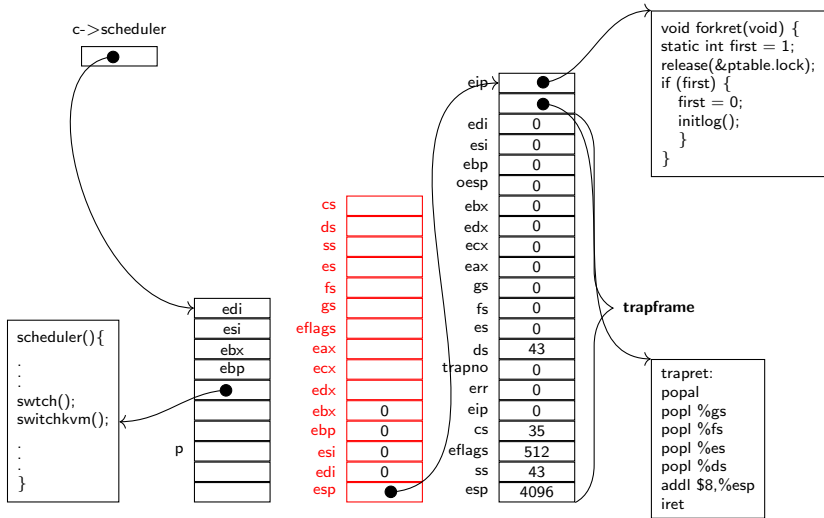
swtch() operation



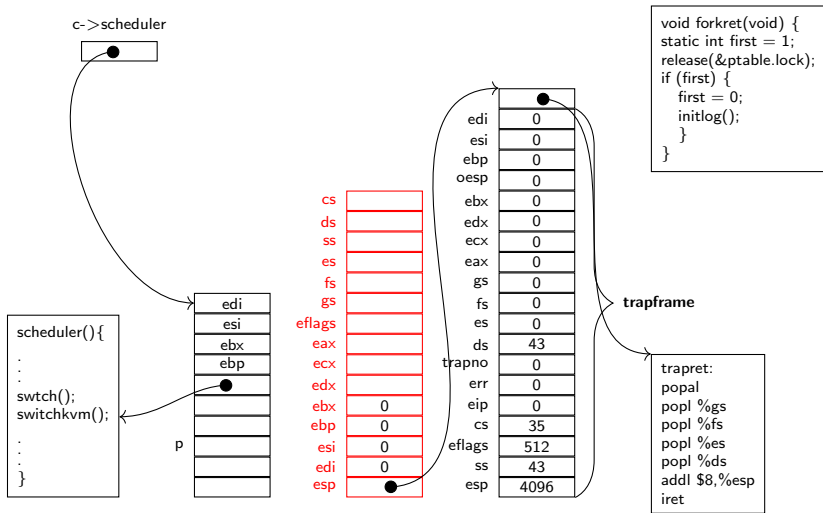
swtch() operation



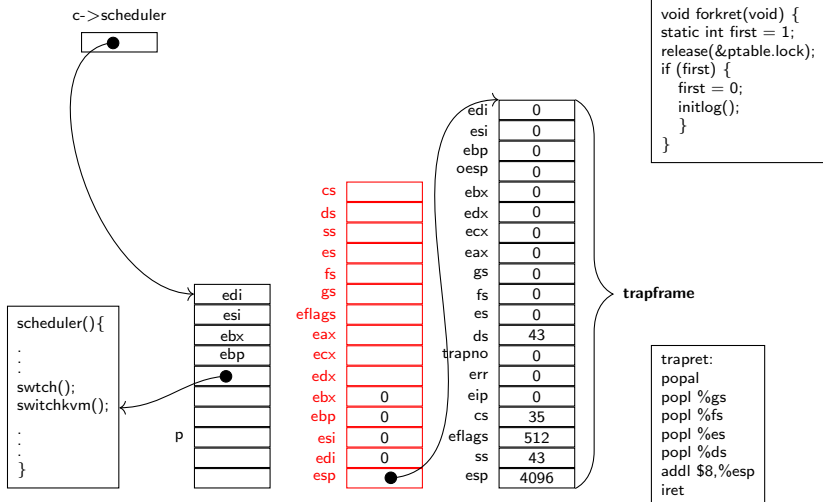
swtch() operation



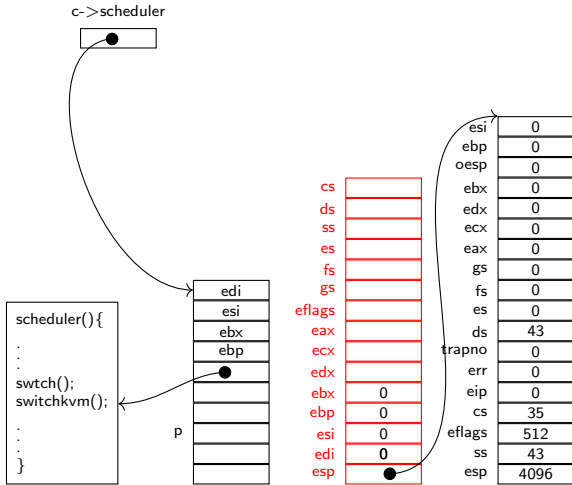
swtch() operation



swtch() operation



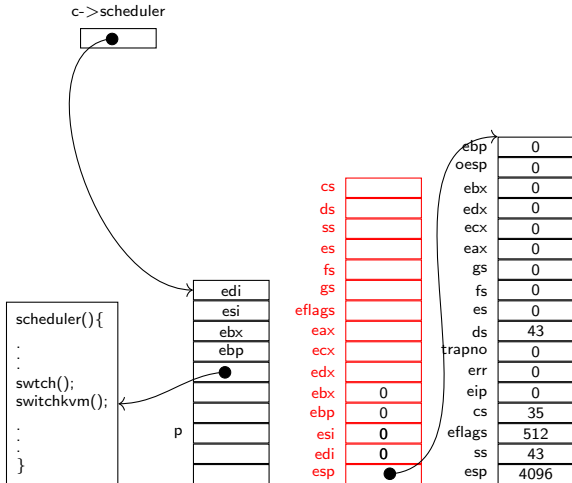
swtch() operation



```
void forkret(void) {
static int first = 1;
release(&ptable.lock);
if (first) {
    first = 0;
    initlog();
}
}
```

```
trapret:
popal
popl %gs
popl %fs
popl %es
popl %ds
add $8,%esp
iret
```

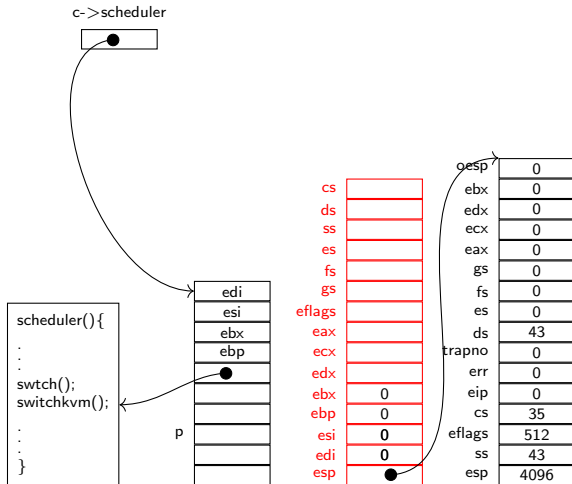
swtch() operation



```
void forkret(void) {
static int first = 1;
release(&ptable.lock);
if (first) {
    first = 0;
    initlog();
}
}
```

```
trapret:
popal
popl %gs
popl %fs
popl %es
popl %ds
addl $8,%esp
iret
```


swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

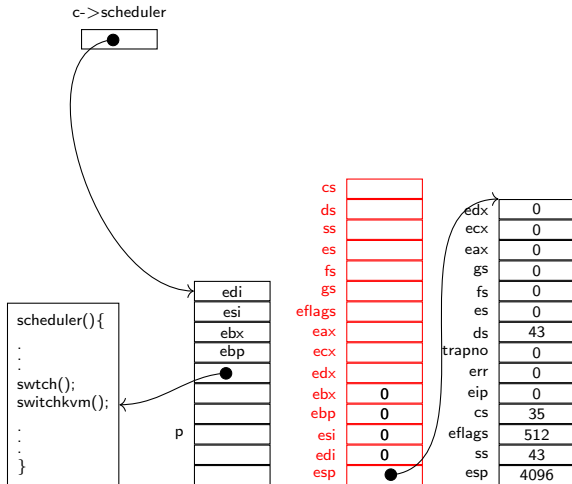
swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    irt
```


swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```



swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

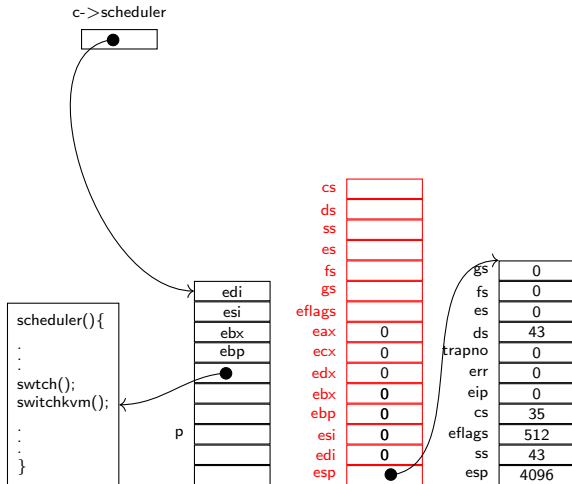
swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    addl $8,%esp  
    iret
```


swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

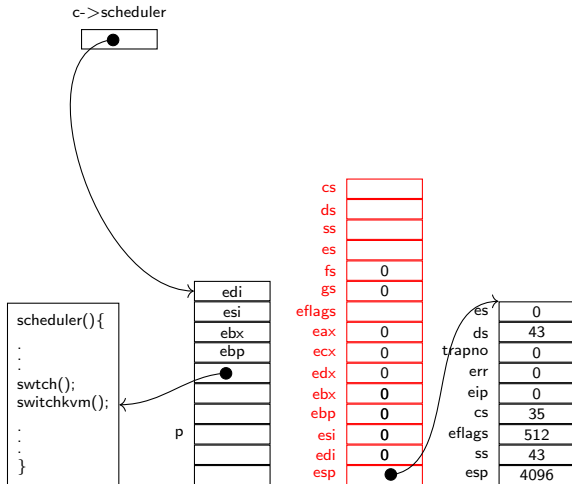
swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

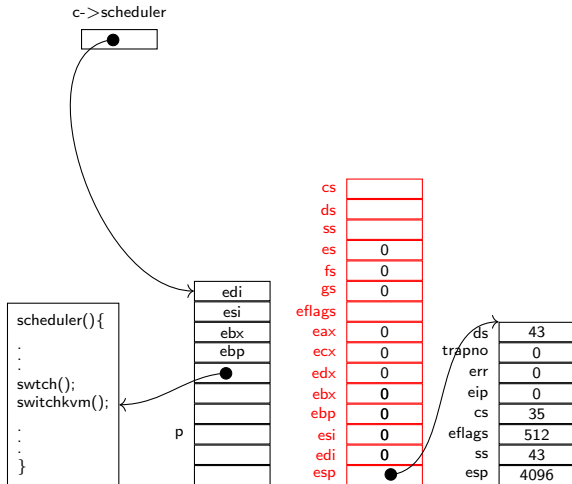
swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```


swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

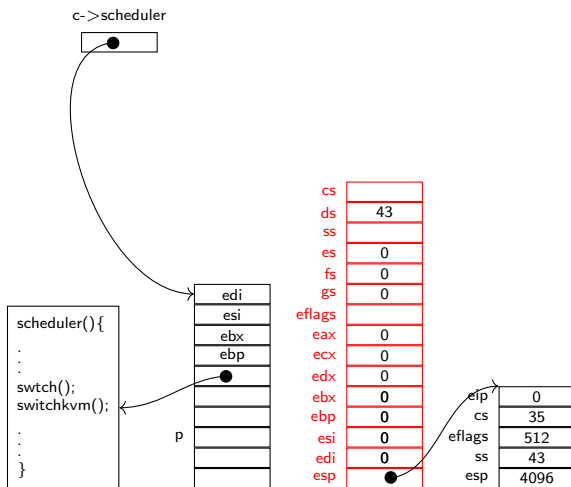
swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

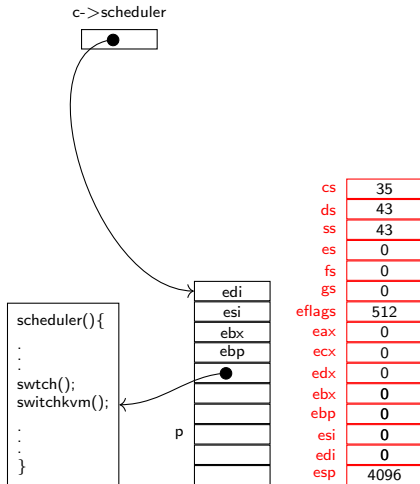
swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```


swtch() operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    irt
```

Context switch

- Calling **switch()**:
 - Creates a **context** structure on the current stack.
 - Stores the **context** structure address created in the first argument.
 - Load the **context** structure pointed to by the second argument.
- We are switching KERNEL contexts.
- User mode context of a process is loaded by the kernel side of the process.

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????
 - The **gcc** calling conventions deals with them!

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????
 - The **gcc** calling conventions deals with them!
- Where are **cs**, **ds**, and **ss**????

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????
 - The **gcc** calling conventions deals with them!
- Where are **cs**, **ds**, and **ss**????
 - The base and limit fields are identical across all kernel sides.
 - So, they need to be loaded only on each kernel entering.

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????
 - The **gcc** calling conventions deals with them!
- Where are **cs**, **ds**, and **ss**????
 - The base and limit fields are identical across all kernel sides.
 - So, they need to be loaded only on each kernel entering.
- Where is **gdt**????

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????
 - The **gcc** calling conventions deals with them!
- Where are **cs**, **ds**, and **ss**????
 - The base and limit fields are identical across all kernel sides.
 - So, they need to be loaded only on each kernel entering.
- Where is **gdtr**????
 - The address and size fields are different across processors.
 - The base and limit **MUST NOT** change between kernel sides on the same CPU.
 - Since **gdtr** is privileged, it needs to be loaded **ONLY** on kernel initialization.

If switching to user mode is expected:

- The `tr` register should contain the index of a TSS descriptor.
- The TSS descriptor should point to a `taskstate` structure.
- The `ss0` and `esp0` fields should point to a valid kernel stack top.
- The above is ESSENTIAL for proper interrupt service in user mode.

switchvm

```
1860 void switchvm(struct proc *p) {
    pushcli();
    mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A,
                                   &mycpu()->ts,
                                   sizeof(mycpu()->ts)-1, 0);
    mycpu()->gdt[SEG_TSS].s = 0;
    mycpu()->ts.ss0 = SEG_KDATA<<3;
    mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;
    mycpu()->ts.iomb = (ushort) 0xFFFF;
    ltr(SEG_TSS << 3);
    if (p->pgdir == 0)
        panic("switchvm: _no_pgdir");
    lcr3(v2p(p->pgdir)); // switch to new address space
    popcli();
}
```

taskstate (hardware structure)

link	
esp0	
	ss0
esp1	
	ss1
esp2	
	ss2
cr3	
eip	
eflags	
eax	
ecx	
edx	
ebx	
esp	
ebp	
esi	
edi	

	es
	cs
	ss
	ds
	fs
	gs
	ldt
	t
iomb	

taskstate in C

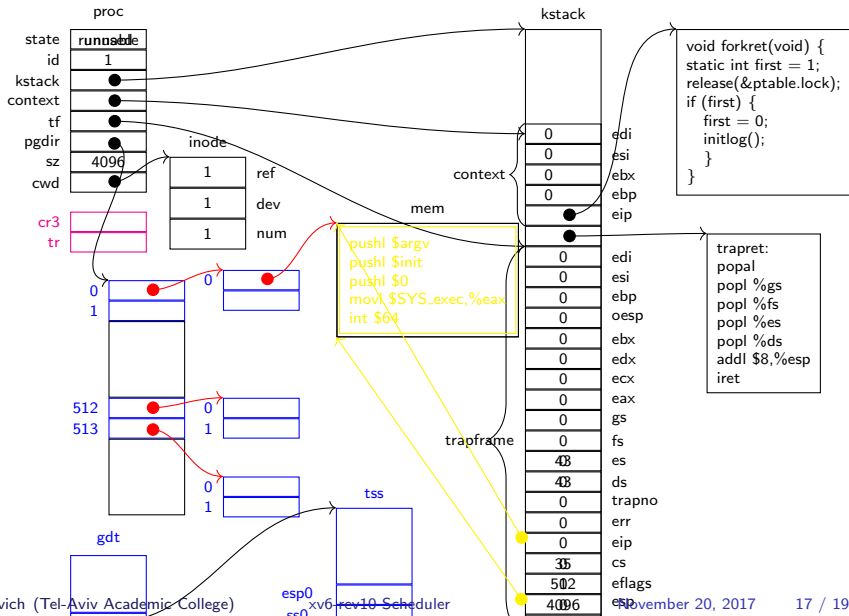
851

```
struct taskstate
{
    uint link;
    uint esp0;
    ushort ss0;
    ushort padding1;
    uint *esp1;
    ushort ss1;
    ushort padding2;
    uint *esp2;
    ushort ss2;
    ushort padding3;
    void *cr3;
    uint *eip;
    uint eflags;
    uint eax;
    uint ecx;
```

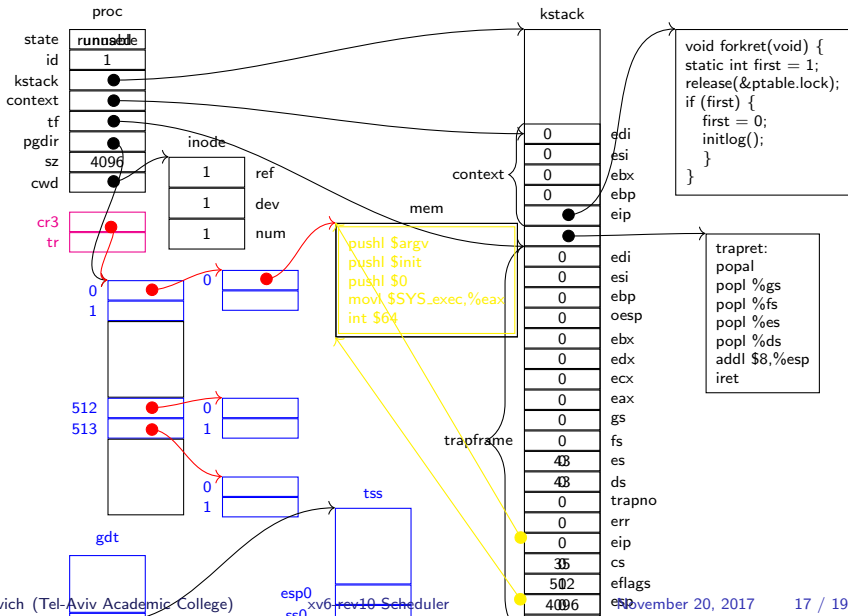
```
    uint edx;
    uint ebx;
    uint *esp;
    uint *ebp;
    uint esi;
    uint edi;
    ushort es;
    ushort padding4;
    ushort cs;
    ushort padding5;
    ushort ss;
    ushort padding6;
    ushort ds;
    ushort padding7;
    ushort fs;
    ushort padding8;
```

```
    ushort gs;
    ushort padding9;
    ushort ldt;
    ushort padding10;
    ushort t;
    ushort iomb;
};
```

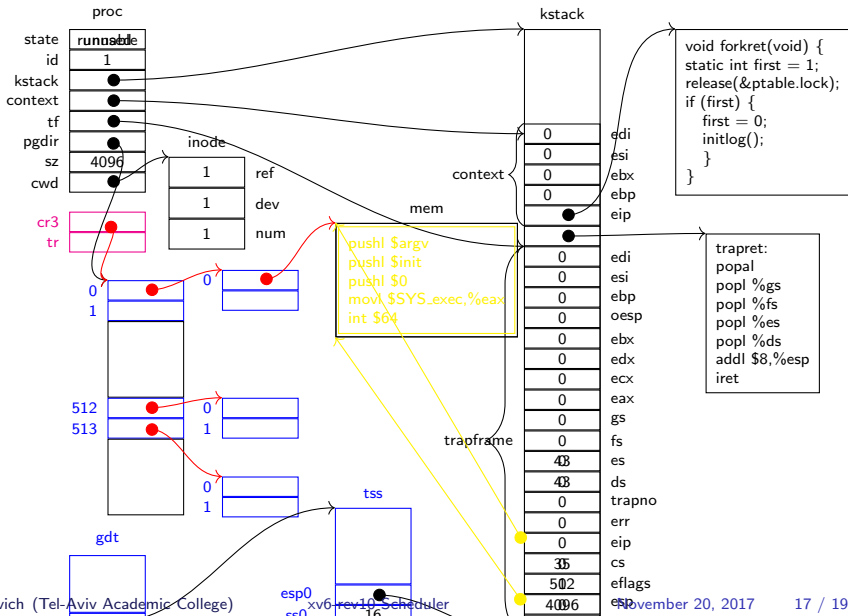
switchvm




switchvm



switchvm



switchvm



co-routines

- The scheduler switches to a process by using:

```
2478      swtch(&c->scheduler , p->context );
```

- A process leaves the cpu by returning to the scheduler using:

```
2516      swtch(&p->context , mycpu()->scheduler );
```

- We have here co-routines.

Event driven kernel

- At this point there is no more LINEAR EXECUTION of the kernel.
- The kernel as of now is EVENT DRIVEN.