

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
Locks

Carmi Merimovich

Tel-Aviv Academic College

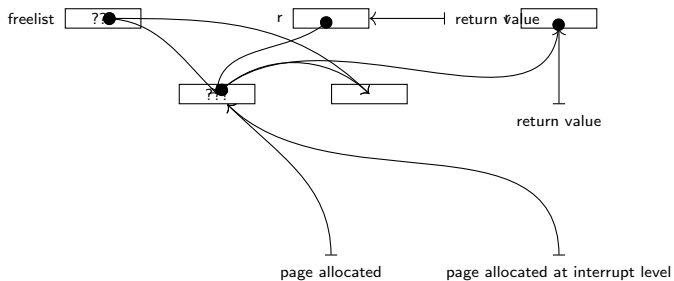
November 30, 2017

What is a uniprocessing risk with kalloc?

```
3187 char *kalloc(void) {  
    struct run *r;  
  
    r = kmem.freelist;  
    if (r)  
        kmem.freelist = r->next;  
  
    return (char*)r;  
}
```

What might happen if kalloc is called within an hardware interrupt routine?

Uniprocessor **kalloc()** risks



Race condition

- When several execution contexts change common resource the behavior might be unpredictable.
- Simple solution:
 - Serialize changes to the common resource.
 - This decreases the amount of parallelism.
 - Hence, performance is decreased.
- Complicated solution:
 - Discover parallel algorithm ...

Race Condition

Uniprocessor

Race condition demonstration

```
extern x=0,y=5;
```

```
x += y;
```

Figure: main code

```
x -= y;
```

Figure: interrupt service code

The compiled code is something like the following:

```
movl x,%eax  
addl y,%eax  
movl %eax,x
```

```
movl x,%eax  
subl y,%eax  
movl %eax,x
```

x value

- Assume the main code executes once.
- Assume the interrupt level code executes once.
- What will be the value of **x** at the end of execution?

Scenario 1

x	8
y	5

eax 0

```
⋮  
⋮  
movl x,%eax  
addl y,%eax  
movl %eax,x  
⋮  
⋮
```

```
⋮  
⋮  
movl x,%eax  
subl y,%eax  
movl %eax,x  
⋮  
⋮
```

x=0

Scenario 2

x	-5
y	5

eax -5

```
⋮  
⋮  
movl x,%eax  
addl y,%eax  
movl %eax,x  
⋮  
⋮
```

```
⋮  
⋮  
movl x,%eax  
subl y,%eax  
movl %eax,x  
⋮  
⋮
```

x=0

Scenario 3

x	5
y	5

eax ~~5~~

```
⋮  
⋮  
movl x,%eax  
addl y,%eax  
movl %eax,x  
⋮  
⋮
```

```
⋮  
⋮  
movl x,%eax  
subl y,%eax  
movl %eax,x  
⋮  
⋮
```

x=5

Hence:

- The above code is NONDETERMINISTIC.
- A simple solution is to (properly) SERIALIZE.

Solution: Masking external interrupts




Figure: eflags

- If `eflags.IF==0` then external interrupts are ignored.
- if `eflags.IF==1` then external interrupts are serviced.

`eflags.IF` can be controlled with the following privileged instructions:

- `cli: eflags.IF \leftarrow 0.`
- `sti: eflags.IF \leftarrow 1.`

Reading whole of `eflags` into (say) `%eax` can be done using:

```
pushfl
popl %eax
```

xv6 routines for controlling eflags.IF

```
557 static inline void cli(void) {  
    asm volatile("cli");  
}  
  
562 static inline void sti(void) {  
    asm volatile("sti");  
}  
  
544 static inline uint readeflags(void) {  
    uint eflags;  
    asm volatile("pushfl;_popl_%0" : "=r" (eflags));  
    return eflags;  
}
```

Serialization

```
extern x=0,y=5;
```

```
cli();
```

```
x += y;
```

```
sti();
```

```
x -= y;
```

Figure: main code

Figure: interrupt service code

- If the interrupt service is not interruptible, the solution is OK.

Serialize at interrupt level

```
extern x=0,y=5;
```

```
cli ();  
x += y;  
sti ();
```

Figure: main code

```
cli ();  
x -= y;  
sti ();
```

Figure: interrupt service code

- If the interrupt service is not interruptible, this solution is not OK.
- Using naked **sti** and **cli** is error prone.

Using cli and sti directly

Problem:

```
func b() {  
    cli();  
    :  
    sti();  
}
```

```
func a() {  
    cli();  
    :  
    b();  
    : // We are in trouble here  
    sti();  
}
```


Solution

We will not call `cli()` and `sti()` directly. Instead:

- We add a global `uint ncli`.
- On each `cli()` needed we call `pushcli()` which will:
 - `cli()`.
 - Increase `ncli`.
- On each `sti()` needed we call `popcli()` which will:
 - Decrease `ncli`.
 - On transition of `ncli` from 1 to 0, call `sti()`.

Solution code

```
int  ncli=0;  // Multiprocessors problem!
```

```
void pushcli(void) {  
    cli();  
    ncli++;  
}
```

```
void popcli(void) {  
    if (--ncli == 0)  
        sti();  
}
```

Uniprocessor interrupt service masking

Problem:

```
func b() {  
    pushcli();  
    :  
    popcli();  
}
```

```
func a() {  
    pushcli();  
    :  
    b();  
    :  
    popcli();  
}
```

eflags

- In the above `pushcli()/popcli()` implementation it was implicit before the outermost `pushcli()` interrupts were enabled.
- This is not necessarily true (e.g., xv6 initialization vs. rest of the code).
- In order to handle this we add the following.
- On transition of `ncli` from 0 to 1 we save `eflags.IF`.
- On transition of `ncli` from 1 to 0 we restore `eflags.IF`.

Uniprocessor solution code

```
#define FL_IF 0x00000200
int ncli=0, intena;  // Multiprocessors problem!

void pushcli(void) {
    int eflags = readeflags();
    cli();
    if (ncli++ == 0)
        intena = eflags & FL_IF;
}

void popcli(void) {
    if (--ncli == 0 && intena)
        sti();
}
```

Multiprocessor solution code

```
#define FL_IF 0x00000200
```

```
1655 void pushcli(void) {  
    int eflags = readeflags();  
    cli();  
    if (mycpu()->ncli == 0)  
        mycpu()->intena = eflags & FL_IF;  
    mycpu()->ncli += 1;  
}  
  
1667 void popcli(void) {  
    if (--mycpu()->ncli == 0 && cpu->intena)  
        sti();  
}
```

xv6 Naked pushcli/popcli

Only in switchvm:

```
1860 void switchvm(struct proc *p) {
    pushcli();
    mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A,
                                &mycpu()->ts,
                                sizeof(mycpu()->ts)-1, 0);
    mycpu()->gdt[SEG_TSS].s = 0;
    mycpu()->ts.ss0 = SEG_KDATA<<3;
    mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;
    mycpu()->ts.iobm = (ushort) 0xFFFF;
    ltr(SEG_TSS << 3);
    if (p->pgdir == 0)
        panic("switchvm: _no_pgdir");
    lcr3(v2p(p->pgdir)); // switch to new address space
    popcli();
}
```

Race condition

Multiprocessors

What is the multiprocessing risk with allocproc?

2480

```
for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
    if (p->state == UNUSED)  
        goto found;  
return 0;
```

found:

```
p->state = EMBRYO;  
p->pid = nextpid++;
```

What is the multiprocessing risk with scheduler?

```
2770  for (p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
      if (p->state != RUNNABLE)  
          continue;
```

Race condition demonstration

```
extern x=0;
```

```
x += 5;
```

Figure: code on processor 0

```
x -= 5;
```

Figure: code on processor 1

The compiled code is something like the following:

```
addl $5, x
```

```
subl $5, x
```

However, the RAM is passive so in reality the CPU microsteps are:

```
%tmp ← x;  
%tmp ← %tmp + 5;  
x ← %tmp;
```

```
%tmp ← x;  
%tmp ← %tmp - 5;  
x ← %tmp;
```

Atomic instruction on MP systems

- x86:

```
lock; xchgl mem, reg
```

- xv6 function:

```
569 static inline uint xchg(volatile uint *addr,
                           uint newval) {
    uint result;

    asm volatile("lock; _xchgl _%0, _%1" :
                  "+m" (*addr), "=a" (result) :
                  "1" (newval) :
                  "cc");
    return result;
}
```

Effect of the xchg function

- The call

```
old = xchg(&lock , 1);
```

achieves

```
old ← lock;
```

```
lock ← 1;
```

with no other processor executing other locked instructions in parallel.

- If we have recursive interrupts using lock then pushcli()/popcli() should be added.

spinlock structure

```
1501 struct spinlock {  
    uint locked; // Is the lock held?  
  
    // For debugging:  
    char *name; // Name of lock.  
    struct cpu *cpu; // The cpu holding the lock.  
    uint pcs[10]; // The call stack (an array of program  
    // that locked the lock.  
};  
  
1562 void initlock(struct spinlock *lk, char *name) {  
    lk->name = name;  
    lk->locked = 0;  
    lk->cpu = 0;  
}
```

```
1574 acquire(struct spinlock *lk) {  
    pushcli(); // disable interrupts to avoid deadlock.  
  
    while(xchg(&lk->locked, 1) != 0);  
}  
  
1601 void release(struct spinlock *lk) {  
    xchg(&lk->locked, 0);  
    popcli();  
}
```

Race condition reason and solution

- Resources shared by recursive interrupts service routines:
 - `pushcli()`.
 - `popcli()`.
- Resources shared by different processors:
 - `acquire()`.
 - `release()`.

Solving potential race condition in `allocproc` and in `scheduler`.

allocproc with the lock

```
2478 acquire(&ptable.lock);  
for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
    if (p->state == UNUSED)  
        goto found;  
return 0;
```

```
found:  
p->state = EMBRYO;  
p->pid = nextpid++;  
release(&ptable.lock);
```

where the defending spinlock is in ptable:

```
2410 struct {  
    struct spinlock lock;  
    struct proc proc[NPROC];  
} ptable;
```

scheduler

```
2758 void scheduler(void) {  
    struct proc *p;  
    for(;;) { sti();  
  
        acquire(&ptable.lock);  
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {  
            if (p->state != RUNNABLE) continue;  
  
            proc = p;  
            switchvm(p);  
            p->state = RUNNING;  
            switch(&cpu->scheduler, proc->context);  
            switchkvm();  
  
            proc = 0;  
        }  
        release(&ptable.lock);  
    }  
}
```

forkret

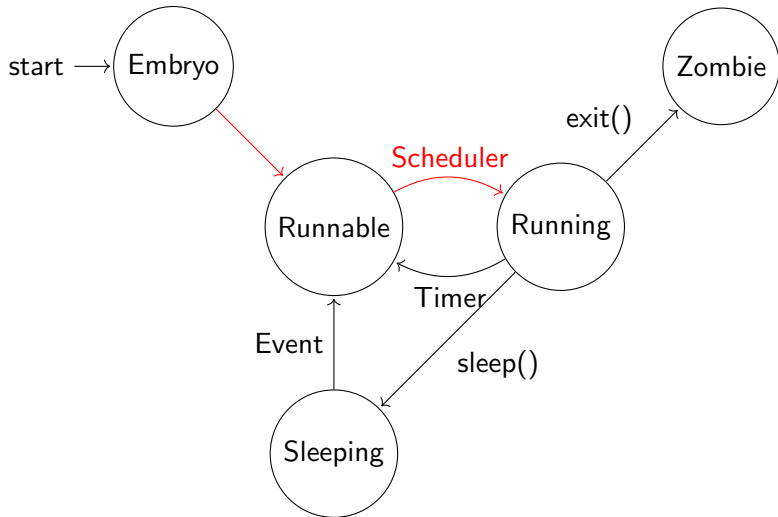
```
2853 forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

lock rules for **swtch()**

- Always call **swtch()** with **ptable.lock** acquired.
- No other lock should be **acquired** on entering **swtch()**.
- This lock will pass on to the scheduler/process, where in due time it will be **relased**.

Never ever return to user mode with **acquired** lock or masked interrupts.

state transition



Going from process to scheduler

Leaving process context

```
2808 void sched(void) {  
    int intena;  
    struct proc *p = myproc();  
  
    intena = mycpu()->intena;  
    swtch(&p->context, mycpu()->scheduler);  
    mycpu()->intena = intena;  
}
```


yielding

```
2828 void yield(void) {  
    acquire(&ptable.lock);  
    myproc()->state = RUNNABLE;  
    sched();  
    release(&ptable.lock);  
}
```

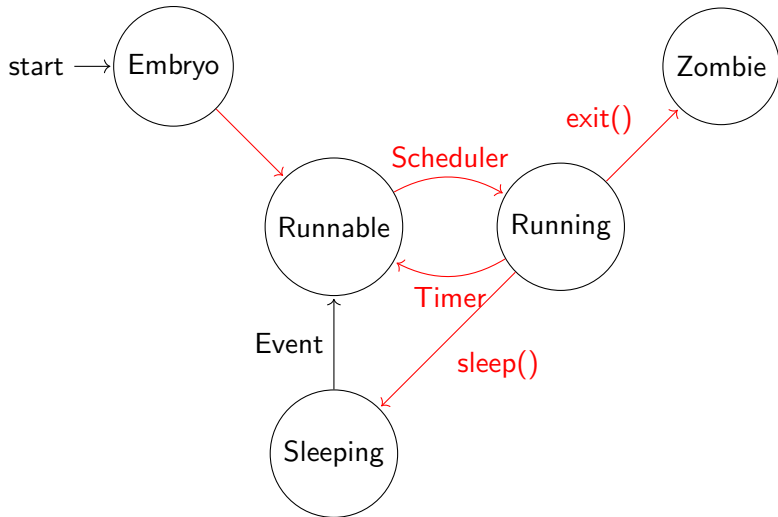
Zombieing

```
2649  acquire(&ptable.lock);  
      :  
2663  myproc()->state = ZOMBIE;  
      sched();
```

sleeping

```
2874 void sleep(void *chan, struct spinlock *lk) {  
    struct proc *p = myproc();  
    if (lk != &ptable.lock) {  
        acquire(&ptable.lock);  
        release(lk);  
    }  
    p->chan = chan;  
    p->state = SLEEPING;  
    sched();  
    p->chan = 0;  
  
    if (lk != &ptable.lock) {  
        release(&ptable.lock);  
        acquire(lk);  
    }  
}
```

state transition



All

transition are due to INTERRUPTS.