

xv6©-rev10  
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)  
syscall arguments

Carmi Merimovich

Tel-Aviv Academic College

January 5, 2017

## Fetching from user mode routines

Safe fetching from user mode using address:

```
int fetchint(int addr, int *i);  
int fetchstr(int addr, char *s);
```

Safe fetching from user mode using argument number:

```
int argint(int n, int *i);  
int argptr(int n, void *p, int size);  
int argstr(int n, char *s);
```

## Fetching a byte

- Assume the definition **uint *addr*** is given.
- How do we fetch, in C, the byte in address ***addr***?

## Fetching a byte

- Assume the definition **uint *addr*** is given.
- How do we fetch, in C, the byte in address ***addr***?
- **`char b = *(char *)addr;`**

## Fetching a byte

- Assume the definition **uint *addr*** is given.
- How do we fetch, in C, the byte in address ***addr***?
- **`char b = *(char *)addr;`**
- Assume we are in the kernel, and ***addr*** was supplied by user mode.
- What is the problem?

## Fetching a byte

- Assume the definition **uint addr**s is given.
- How do we fetch, in C, the byte in address **addr**s?
- **char b = \*(char \*)addr**s;
- Assume we are in the kernel, and **addr**s was supplied by user mode.
- What is the problem?
  - **addr**s is not trustable.

## Fetching a byte

- Assume the definition **uint addr**s is given.
- How do we fetch, in C, the byte in address **addr**s?
- **char b = \*(char \*)addr**s;
- Assume we are in the kernel, and **addr**s was supplied by user mode.
- What is the problem?
  - **addr**s is not trustable.
- What do we do?

## Fetching a byte

- Assume the definition **uint addr**s is given.
- How do we fetch, in C, the byte in address **addr**s?
- **char b = \*(char \*)addr**s;
- Assume we are in the kernel, and **addr**s was supplied by user mode.
- What is the problem?
  - **addr**s is not trustable.
- What do we do?

```
if (addr < proc->sz)
    b = *(char *)addr;
else
    //
    //  handle the error
    //
```



## fetchbyte

```
int fetchbyte(uint addr, char *cp) {  
    if (addr >= proc->sz)  
        return (-1);  
    *cp = *(char *)addr;  
    return (0);  
}
```

## fetchbyte

```
int fetchbyte(uint addr, char *cp) {  
    if (addr >= proc->sz)  
        return (-1);  
    *cp = *(char *)addr;  
    return (0);  
}
```

- What changes should be done in order to get fetchshort?

## fetchshort

```
int fetchshort(uint addr, short *sp) {  
    if (addr >= proc->sz || addr+1 >= proc->sz)  
        return (-1);  
    *sp = *(short *)addr;  
    return (0);  
}
```

## fetchshort

```
int fetchshort(uint addr, short *sp) {  
    if (addr >= proc->sz || addr+1 >= proc->sz)  
        return (-1);  
    *sp = *(short *)addr;  
    return (0);  
}
```

- Do we really need both comparisons above?

## fetchshort

```
int fetchshort(uint addr, short *sp) {  
    if (addr >= proc->sz || addr+1 >= proc->sz)  
        return (-1);  
    *sp = *(short *)addr;  
    return (0);  
}
```

- Do we really need both comparisons above?
  - Yes.


## Fetching **long** from an untrusted user address

- Fetching a **long** from address **addr** means fetching four bytes from addresses **addr**, **addr+1**, **addr+2**, and **addr+3**.
- All these addresses should all be below **proc->sz**.
- The **fetchint** routine returns -1 if the above is not correct.
- If it is correct, the content of the **long** is put into **\*ip**., and the functions returns **0**.


3567

```
int fetchint(uint addr, int *ip) {  
    if (addr >= myproc()->sz || addr+4 > myproc()->sz)  
        return -1;  
    *ip = *(int*)(addr);  
    return 0;  
}
```

## One argument syscall




## One argument syscall







## One argument syscall



## One argument syscall



## One argument syscall




**int \$64**

# One argument syscall


esp	??
eax	??



# One argument syscall




## One argument syscall



**\*myproc()->tf->esp**


## One argument syscall



**$*(myproc() \rightarrow tf \rightarrow esp + 4)$**


**This is the argument!**

## Two arguments syscall







## Two arguments syscall




## Two arguments syscall




## Two arguments syscall



## Two arguments syscall



## Two arguments syscall




**int \$64**

## Two arguments syscall


esp	??
eax	??



## Two arguments syscall




## Two arguments syscall



**\*myproc()->tf->esp**




## Two arguments syscall



**`*(myproc()->tf->esp +4)`**

**This is argument 0!**

## Two arguments syscall



**`*(myproc()->tf->esp + 8)`**  
**This is argument 1!**

## system call args in kernel mode

- The kernel stack replaced the user stack.
- Hence the arguments are NOT on the kernel stack.
- However, the user stack address was saved on the trapframe.
- Thus in the `syscall` routine we have:
  - arg 0 address is `myproc()->tf->esp + 4`.
  - arg 1 address is `proc->tf->esp + 8`.
  - $\vdots$
  - arg n address is `proc->tf->esp + 4 + 4 \times n`.
- `proc->tf->esp` was really set by the user, hence untrusted.
- Any computation derived from `proc->tf->esp` is also untrusted.


## Fetching the **n**-th integer argument

```
3602 int argint(int n, int *ip) {  
    return fetchint(myproc()->tf->esp + 4 + 4*n, ip);  
}
```


## Another form for **argint**

```
struct stackform {  
    int ignore;  
    int arg[1]; //new compilers allow arg[]  
};  
  
int argint(int n, int *ip) {  
    return fetchint(&proc->tf->esp->arg[n], ip);  
}
```


## Buffer argument



## Buffer argument




## Buffer argument






## Buffer argument




## Buffer argument




**int \$64**

## Buffer argument


esp    ??  
eax    ??



## Buffer argument



## Buffer argument



```
addr=*(proc->tf->esp +4)
      *(addr+i)
```

## Buffer argument

```
3303 int argptr(int n, char **pp, int size) {  
    int addr;  
  
    if (argint(n, &addr) < 0)  
        return -1;  
    if ((uint)addr >= proc->sz ||  
        (uint)addr+size > proc->sz)  
        return -1;  
    *pp = (char*)i;  
    return 0;  
}
```

## String argument

```
3320 int argstr(int n, char **pp) {  
    int addr;  
    if (argint(n, &addr) < 0)  
        return -1;  
    return fetchstr(addr, pp);  
}  
  
3278 int fetchstr(uint addr, char **pp) {  
    char *s, *ep;  
  
    if (addr >= proc->sz) return -1;  
    *pp = (char*)addr;  
    ep = (char*)proc->sz;  
    for (s = *pp; s < ep; s++)  
        if (*s == 0) return s - *pp;  
    return -1;  
}
```