

# The inode Layer II

xv6-rev7

Carmi Merimovich

CS School, Tel Aviv Academic College

January 2, 2017

## (memory) inode

```
3671 #define NDIRECT 12

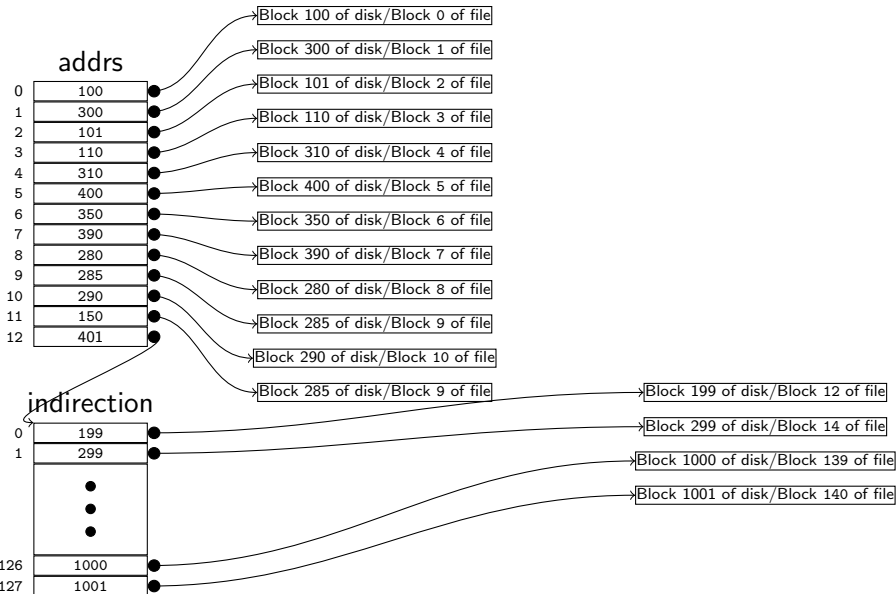
3762 struct inode {
    uint dev; // Device number
    uint inum; // Inode number
    int ref; // Reference count
    int flags; // I_BUSY, I_VALID

    short type; // copy of disk inode
    short major;
    short minor;
    short nlink;
    uint size;
    uint addrs[NDIRECT+1];
};

#define I_BUSY 0x1
#define I_VALID 0x2
```

We aim at readi.  
Beforehand we need bmap for which we need addrs.

## addrs



## bmap(ip,bn)

- Return the physical block # which block #bn of the file ip uses.
- If asked for a non existent block, allocates it.
- Allocation uses the buffer layer function balloc().
- If the inode is updated it is the caller responsibility to call iupdate.

## bmap() (1)

```
4810 bmap(struct inode *ip, uint bn) {  
    uint addr, *a;  
    struct buf *bp;  
  
    if (bn < NDIRECT) {  
        if ((addr = ip->addrs[bn]) == 0)  
            ip->addrs[bn] = addr = balloc(ip->dev);  
        return addr;  
    }  
}
```

## bmap() (2)

```
bn -= NDIRECT;
```

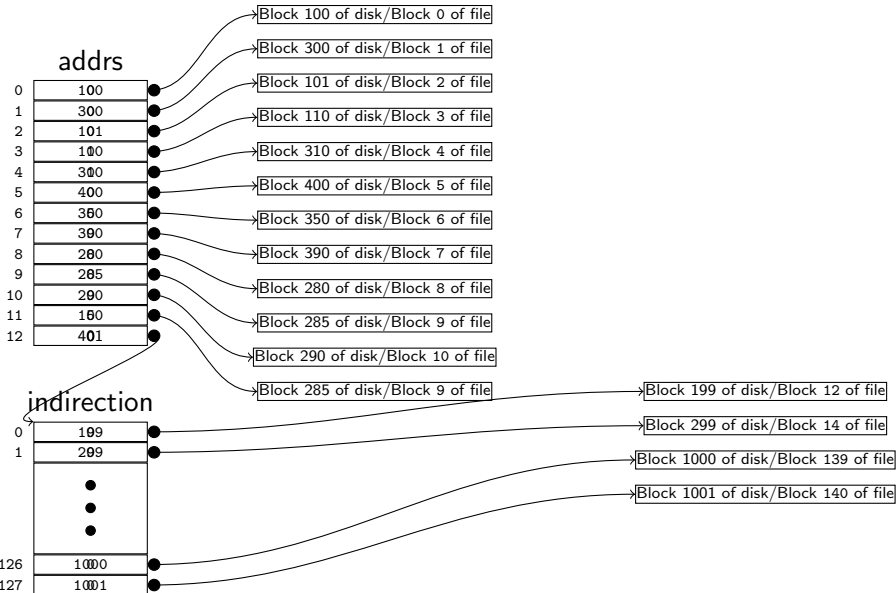
```
if (bn < NINDIRECT) {  
    if ((addr = ip->addrs[NDIRECT]) == 0)  
        ip->addrs[NDIRECT] = addr = balloc(ip->dev);  
    bp = bread(ip->dev, addr);  
    a = (uint*)bp->data;  
    if ((addr = a[bn]) == 0) {  
        a[bn] = addr = balloc(ip->dev);  
        log_write(bp);  
    }  
    brelse(bp);  
    return addr;  
}  
panic("bmap: out of range");  
}
```

## readi

```
4902 readi(struct inode *ip, char *dst, uint off, uint n) {
    uint tot, m;
    struct buf *bp;
    if (ip->type == T_DEV) {
        if (ip->major < 0 || ip->major >= NDEV ||
            !devsw[ip->major].read)
            return -1;
        return devsw[ip->major].read(ip, dst, n);
    }
    if (off > ip->size || off + n < off) return -1;
    if (off + n > ip->size)
        n = ip->size - off;
    for (tot=0; tot<n; tot+=m, off+=m, dst+=m) {
        bp = bread(ip->dev, bmap(ip, off/BSIZE));
        m = min(n - tot, BSIZE - off%BSIZE);
        memmove(dst, bp->data + off%BSIZE, m);
        brelse(bp);
    }
    return n;
}
```



## Filling in addr



## writei (1)

```
4950 int writei(struct inode *ip, char*src, uint off, uint n,
          uint tot, m;
          struct buf *bp;

if (ip->type == T_DEV) {
    if (ip->major < 0 || ip->major >= NDEV || !devsw[ip->major])
        return -1;
    return devsw[ip->major].write(ip, src, n);
}

if (off > ip->size || off + n < off) return -1;
if (off + n > MAXFILE*BSIZE) return -1;
```

## writei (2)

```
for (tot=0; tot<n; tot+=m, off+=m, src+=m) {  
    bp = bread(ip->dev, bmap(ip, off/BSIZE));  
    m = min(n - tot, BSIZE - off%BSIZE);  
    memmove(bp->data + off%BSIZE, src, m);  
    log_write(bp);  
    brelse(bp);  
}  
  
if (n > 0 && off > ip->size) {  
    ip->size = off;  
    iupdate(ip);  
}  
return n;  
}
```

## itrunc (1)

```
4855 static void itrunc(struct inode *ip) {  
    int i, j;  
    struct buf *bp;  
    uint *a;  
  
    for (i = 0; i < NDIRECT; i++){  
        if (ip->addrs[i]){  
            bfree(ip->dev, ip->addrs[i]);  
            ip->addrs[i] = 0;  
        }  
    }  
}
```

## itrunc (2)

```
if (ip->addrs[NDIRECT]) {
    bp = bread(ip->dev, ip->addrs[NDIRECT]);
    a = (uint*)bp->data;
    for(j = 0; j < NINDIRECT; j++) {
        if(a[j])
            bfree(ip->dev, a[j]);
    }
    brelse(bp);
    bfree(ip->dev, ip->addrs[NDIRECT]);
    ip->addrs[NDIRECT] = 0;
}

ip->size = 0;
iupdate(ip);
}
```

# XV6 on disk structure: The big picture

