

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
The **fork** system call

Carmi Merimovich

Tel-Aviv Academic College

November 30, 2017

sys_fork

```
3760 int sys_fork(void) {  
    return fork();  
}
```

fork()


Recall:

- A child (of the invoker) process is created.
- The `pid` of the child process is returned to the invoker.
- The child process is (almost) identical to the parent process.
 - To the child, the return value of the system call is zero.
- So, how do we begin?

Child process state needed

eax	0
ebx	pebx
ecx	pecx
edx	pedx
ebp	pebp
esi	pesi
edi	pedi
esp	pesp
eip	peip

cs	pcs
ds	pds
ss	pss
es	pes
fs	pfs
gs	pgs



proc struct

How do we fill the fields of the new process?

```
uint sz; // @proc->sz@
pde_t* pgdir; // @Serious replication needed@
char *kstack; // @probably allocproc()@
enum procstate state; // @RUNNABLE@
volatile int pid; // @allocproc()@
struct proc *parent; // @proc@
struct trapframe *tf; // @allocproc()@
struct context *context; // @allocproc()@
void *chan; // @00@
int killed; // @00@
struct file *ofile[NOFILE]; // @filedup()@ (when st
struct inode *cwd; // @idup()@ (when studying fs)
char name[16]; // @proc->name@
};
```

Needed work


- `proc struct` and friends.
- Filling `pgdir` requires considerable code replication.
- Replicatin of `ofile` and `cwd` requires help from the relevant modules.

proc struct and friends

allocproc(): (1) Finding unused proc structure

```
2473 static struct proc *allocproc(void) {  
    struct proc *p;  
    char *sp;  
  
    acquire(&ptable.lock);  
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
        if (p->state == UNUSED)  
            goto found;  
    release(&ptable.lock);  
    return 0;  
  
found:  
    p->state = EMBRYO;  
    p->pid = nextpid++;  
    release(&ptable.lock);
```


allocproc(): (1) Operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```


```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

allocproc: (2) Initialize process kernel stack

2494

```
if ((p->kstack = kalloc()) == 0) {  
    p->state = UNUSED;  
    return 0;  
}  
sp      = p->kstack + KSTACKSIZE;  
sp      -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
sp      -= 4;  
*(uint*)sp = (uint)trapret;  
sp      -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;  
return p;  
}
```

allocproc: (2) Operation



User space replication

(two methods)

Replication page 0 of current process only

Replicating current process first page

- Allocating new block of memory:

```
dst = kalloc();
```

- Copying. One of the following is possible:

1. `memmove(dst, 0, 4096);`

```
p = walkpgdir(myproc()->pgdir, 0, 0);  
memmove(dst, p2v(PTE_ADDR(*p)), 4096);
```

2. This is of course useless as it is.

Create new address space

Replication page 0 of current process only

Add mapping rules in the new address space to the replication

New address space and mapping

- Creating a new address space:

```
pgdir = setupkvm();
```

- Allocate and copy:

```
dst = kalloc();  
p = walkpgdir(myproc()->pgdir, 0, 0);  
memmove(dst, p2v(PTE_ADDR(*p)), 4096);
```

- Adding translation rule:

```
mappages(pgdir, 0, 4096, v2p(dst), (*p) & 4095);
```

Replicating means doing the copy and translation for each page.

Replicate and Map ALL user space pages

Replicating currnet process pages

NO ERROR CHECKING IN HERE!

```
pgdir = setupkvm();  
  
for (va=0; va<myproc()->sz; va += PGSIZE) {  
    kva = kalloc();  
    memmov(kva, va, PGSIZE);  
    pte = walkpgdir(myproc()->pgdir, va, 0);  
    mappages(pgdir, va, PGSIZE, v2p(kva), (*pte) & 4095);  
}
```

- If there is allocation error, all previous allocations must be freed!
- We show freeing on the next slide.

Freeing address space

```
for (i = 0; i < 512; i++) {  
    if ((pgdir[i] & PTE_P) == 0)  
        continue;  
    pgtbl = p2v(pgdir[i] & ~4096);  
    for (j=0; j < 1024; j++) {  
        if (pgtbl[j] & PTE_P) {  
            kfree(p2v(pgtbl[j] & ~4095));  
        }  
    }  
    kfree(pgtbl);  
}  
for (i = 512; i < 1023; i++) {  
    if ((pgdir[i] & PTE_P) == 0)  
        continue;  
    pgtbl = p2v(pgdir[i] & ~4096);  
    kfree(pgtbl);  
}  
kfree(pgdir);
```

xv6 code for user space replication

(More general than needed)

xv6 replicating and freeing address space

The following xv6 code replicates arbitrary address space.

- The code checks for allocation errors.
- It deallocates all previous allocation in case of failure.

copyuvm()

2035

```
pde_t* copyuvm(pde_t *pgdir, uint sz) {
    pde_t *d; pte_t *pte;
    uint pa, i;
    char *mem;
    if ((d = setupkvm()) == 0) return 0;
    for (i = 0; i < sz; i += PGSIZE) {
        if ((pte = walkpgdir(pgdir, (void *) i, 0)) == 0) panic
        if (!(*pte & PTE_P)) panic("copyuvm: _page_not_present")
        pa = PTE_ADDR(*pte);
        if ((mem = kalloc()) == 0) goto bad;
        memmove(mem, (char*)p2v(pa), PGSIZE);
        if (mappages(d, (void*)i, PGSIZE, v2p(mem),
                     PTE_FLAGS(*pte)) < 0) goto bad;
    }
    return d;
bad:
    freevm(d);
    return 0;
}
```

freevm()

```
2003 void freevm(pde_t *pgdir) {
    uint i;

    if (pgdir == 0)
        panic("freevm: _no_pgdir");
    deallocvm(pgdir, KERNBASE, 0);
    for (i = 0; i < NPENTRIES; i++) {
        if (pgdir[i] & PTE_P) {
            char * v = p2v(PTE_ADDR(pgdir[i]));
            kfree(v);
        }
    }
    kfree((char*)pgdir);
}
```

deallocvm

```
1961 deallocvm(pde_t *pgdir, uint oldsz, uint newsz) {  
    pte_t *pte;  
    uint a, pa;  
    if (newsz >= oldsz) return oldsz;  
    a = PGROUNDUP(newsz);  
    for (; a < oldsz; a += PGSIZE) {  
        pte = walkpgdir(pgdir, (char*)a, 0);  
        if (!pte) a += (NPTENTRIES - 1) * PGSIZE;  
        else if ((*pte & PTE_P) != 0) {  
            pa = PTE_ADDR(*pte);  
            if (pa == 0) panic("kfree");  
            char *v = p2v(pa);  
            kfree(v);  
            *pte = 0;  
        }  
    }  
    return newsz;  
}
```


fork

- `allocproc`:
 - An `EMBRYO` `proc` struct is constructed.
 - A kernel stack is allocated.
 - An uninitialized `trapframe` is allocated.
 - An artificial context is constructed.
- The user space memory of the caller is replicated.
- A new matching page table is constructed.
- The `trapframe` of the caller is copied to the uninitialized `trapframe`.
- The `eax` field of the new `trapframe` is cleared.
- File pointers are replicated.
- Rest of the caller **`proc`** struct fields are copied to the new **`proc`** struct.

fork (1)

```
2580 int fork(void) {  
    int i, pid;  
    struct proc *np;  
  
    if ((np = allocproc()) == 0)  
        return -1;  
  
    if ((np->pgdir=copyuvm(myproc()->pgdir, myproc()->sz  
                          == 0) {  
        kfree(np->kstack);  
        np->kstack = 0;  
        np->state = UNUSED;  
        return -1;  
    }  
}
```

fork (2)

```
np->sz = myproc()->sz;  
np->parent = myproc();  
*np->tf = *myproc()->tf;
```

```
np->tf->eax = 0;
```

```
for (i = 0; i < NOFILE; i++)  
    if (myproc()->ofile[i])  
        np->ofile[i] = filedup(myproc()->ofile[i]);  
np->cwd = idup(myproc()->cwd);
```

```
safestrcpy(np->name, myproc()->name, sizeof(myproc()  
pid = np->pid;  
np->state = RUNNABLE;  
return pid;  
}
```

How and when the child runs?!

Recall the scheduler

scheduler

2758

```
void scheduler(void) {  
    struct proc *p;  
    struct cpu *c = mycpu();  
    c->proc = 0;  
    for(;;) { sti();  
        acquire(&ptable.lock);  
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++) {  
            if (p->state != RUNNABLE) continue;  
  
            c->proc = p;  
            switchvm(p);  
            p->state = RUNNING;  
            swch(&c->scheduler, p->context);  
            switchkvm();  
            c->proc = 0;  
        }  
        release(&ptable.lock);  
    }  
}
```

switchvm

If switching to user mode is expected:

- The `tr` register should contain the index of a TSS descriptor.
- The TSS descriptor should point to a `taskstate` structure.
- The `ss0` and `esp0` fields should point to a valid kernel stack top.
- The above is **ESSENTIAL** for proper interrupt service in user mode.

switchvm

```
1860 void switchvm(struct proc *p) {
    pushcli();
    mycpu()->gdt[SEG_TSS] = SEG16(STS_T32A,
                                &mycpu()->ts,
                                sizeof(mycpu()->ts)-1, 0);
    mycpu()->gdt[SEG_TSS].s = 0;
    mycpu()->ts.ss0 = SEG_KDATA<<3;
    mycpu()->ts.esp0 = (uint)p->kstack + KSTACKSIZE;
    mycpu()->ts.iomb = (ushort) 0xFFFF;
    ltr(SEG_TSS << 3);
    if (p->pgdir == 0)
        panic("switchvm: _no_pgdir");
    lcr3(v2p(p->pgdir)); // switch to new address space
    popcli();
}
```


taskstate (hardware structure)

link	
esp0	
	ss0
esp1	
	ss1
esp2	
	ss2
cr3	
eip	
eflags	
eax	
ecx	
edx	
ebx	
esp	
ebp	
esi	
edi	

	es
	cs
	ss
	ds
	fs
	gs
	ldt
	t
iomb	

taskstate in C

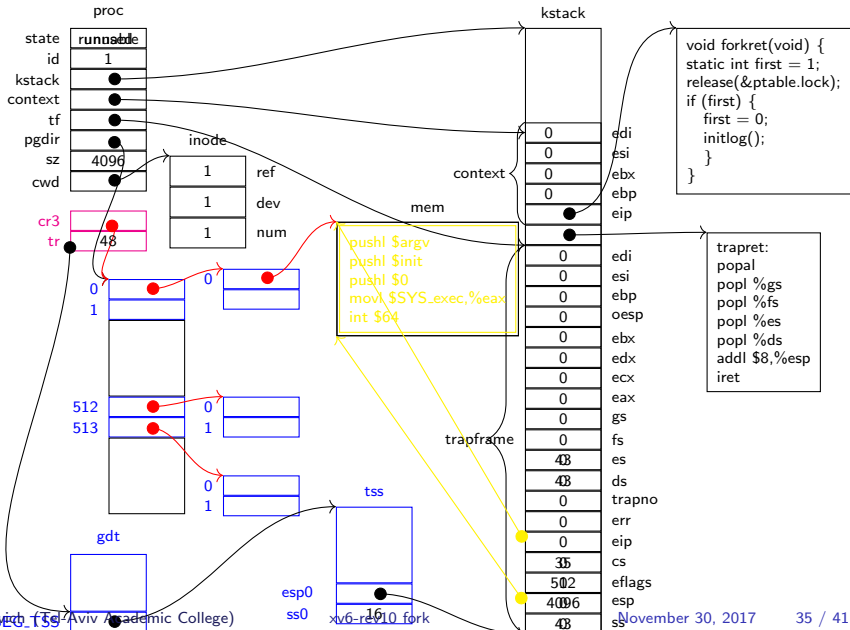
851

```
struct taskstate
{
    uint link;
    uint esp0;
    ushort ss0;
    ushort padding1;
    uint *esp1;
    ushort ss1;
    ushort padding2;
    uint *esp2;
    ushort ss2;
    ushort padding3;
    void *cr3;
    uint *eip;
    uint eflags;
    uint eax;
    uint ecx;
```

```
    uint edx;
    uint ebx;
    uint *esp;
    uint *ebp;
    uint esi;
    uint edi;
    ushort es;
    ushort padding4;
    ushort cs;
    ushort padding5;
    ushort ss;
    ushort padding6;
    ushort ds;
    ushort padding7;
    ushort fs;
    ushort padding8;
```

```
    ushort gs;
    ushort padding9;
    ushort ldt;
    ushort padding10;
    ushort t;
    ushort iomb;
};
```

switchvm



swtch

co-routines

- The scheduler switches to a process by using:

2478 `swtch(&c->scheduler , p->context);`

- A process leaves the cpu by returning to the scheduler using:

2516 `swtch(&p->context , mycpu()->scheduler);`

- We have here co-routines.

swtch()

3058

```
.globl swtch
swtch:
    movl 4(%esp), %eax
    movl 8(%esp), %edx

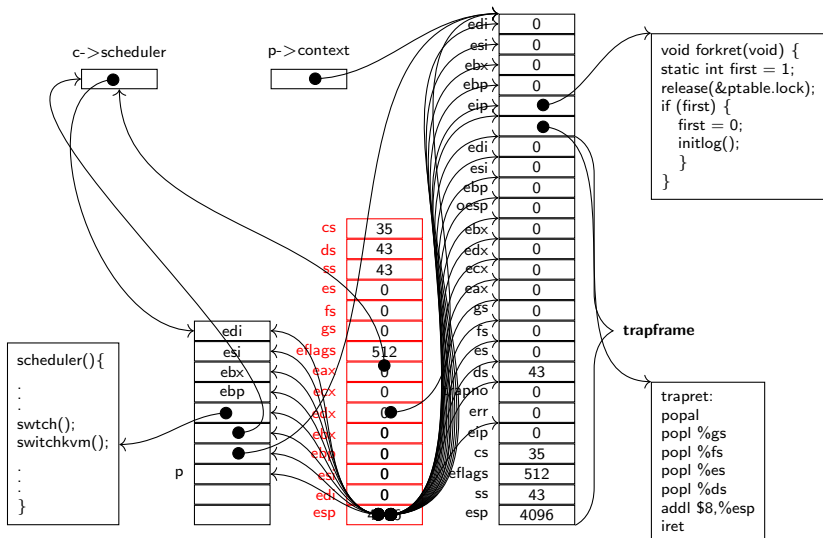
    pushl %ebp
    pushl %ebx
    pushl %esi
    pushl %edi

    movl %esp, (%eax)
    movl %edx, %esp
```

```
    popl %edi
    popl %esi
    popl %ebx
    popl %ebp
    ret
```

The eip field of context is generated by the instruction calling swtch.

swtch() operation



Context switch

- Calling **swtch()**:
 - Creates a **context** structure on the current stack.
 - Stores the **context** structure address created in the first argument.
 - Load the **context** structure pointed to by the second argument.
- We are switching KERNEL contexts.
- User mode context of a process is loaded by the kernel side of the process.

Kernel context seems too small

- Where are **eax**, **ecx**, **edx**????
 - The **gcc** calling conventions deals with them!
- Where are **cs**, **ds**, and **ss**????
 - The base and limit fields are identical across all kernel sides.
 - So, they need to be loaded only on each kernel entering.
- Where is **gdtr**????
 - The address and size fields are different across processors.
 - The base and limit **MUST NOT** change between kernel sides on the same CPU.
 - Since **gdtr** is privileged, it needs to be loaded **ONLY** on kernel initialization.