

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
Entering the Kernel, main

Carmi Merimovich

Tel-Aviv Academic College

September 18, 2017

Where are we?

1. POWERUP.

- The Boot processor executes instructions. MMU inactive.

2. ROM code loads 512-bytes boot block to address 0x07C0 and up.

- ROM terminates by JMPing to address 0x07C0.

3. Boot block code loads the kernel from ide0 into 0x00100000 and up.

- Boot block code terminates by JMPing into the kernel's entry point.
- The entry point address is found at a fixed location in the kernel ELF.
- The entry point code is in Assembly and is labeled entry.

4. entry sets up a temporary kernel programming model.

- MMU is activated with a table coding the following translation:

$$[0x80000000, 0x803FFFFFF] \mapsto [0x00000000, 0x003FFFFFF]$$

- esp is set to the end of a 4KB buffer.
- entry finishes by JMPing into main.

main

Since we got here from **entry**, this code is run by the Boot processor.

- **main** calls routines initializing the different subsystems.
- We study each initialization together with its subsystem.
- Most subsystems need one initialization per (computer) system.
- Initialization which are needed per-cpu are done in **mpmain**.

The application processors begins at **entryother**, then proceed to **mpenter**.

main

1216

```
int main(void) {  
    kinit1(end,  
           P2V(4*1024*1024));  
    kvmalloc();  
    mpinit();  
    lapicinit();  
    seginit();  
    picinit();  
    ioapicinit();  
    consoleinit();  
    uartinit();  
    pinit();  
    tvinit();  
    binit();  
}
```

1251

```
    fileinit();  
    iinit();  
    ideinit();  
    if (!ismp) timerinit();  
    startothers();  
    kinit2(P2V(4*1024*1024),  
           P2V(PHYSTOP));  
    userinit();  
    mpmain();  
}
```

```
static void mpmain(void) {  
    idtinit();  
    xchg(&(mycpu()->started),  
         scheduler());  
}
```

Starting auxiliary processors

In **startothers()**, for each application processor the boot processor executes the following:

1. Copies the **entryother** code into address 0x80007000.
 - **entryother** is a replacement of the **entry** routine.
2. Through the lapic instructs the AP to start executing at 0x0700.
3. Waits for the AP be done with initialization.

Application processor initialization

1. entryother sets up a temporary kernel programming model.
 - MMU is activated with a table coding the following translation:

$[0x80000000, 0x803FFFFFF] \mapsto [0x00000000, 0x003FFFFFF]$

- esp is set to the end of a 4KB buffer.
- entryother finishes by JMPing into mpenter.

```
1240 static void mpenter(void) {  
    switchkvm();  
    seginit();  
    lapicinit();  
    mpmain();  
}
```

The initializations we study now

```
kinit1(end, P2V(4*1024*1024)); // phys page allocation
kvmalloc(); // kernel page table
:
segininit(); // set up segments
:
pinit(); // process table
:
:
kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must copy
userinit(); // first user process
mpmain();
```