

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
First Process Creation, II

Carmi Merimovich

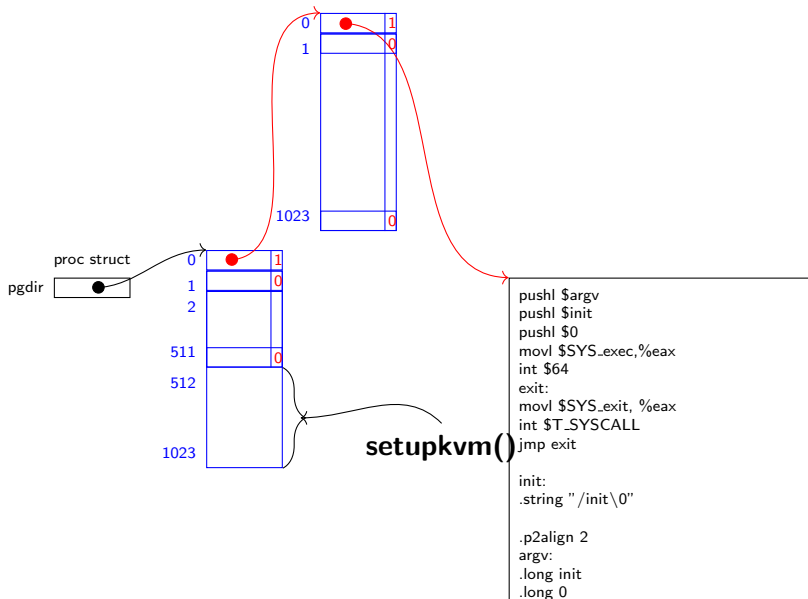
Tel-Aviv Academic College

November 20, 2017

Context

```
kinit1(end, P2V(4*1024*1024)); // phys page allocation
kvmalloc(); // kernel page table
:
segininit(); // set up segments
:
pinit(); // process table
:
:
kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must copy
userinit(); // first user process
mpmain();
```

Context in `userinit()`



The Road to User Mode

Initial user mode state for first process

eax	0
ebx	0
ecx	0
edx	0
ebp	0
esi	0
edi	0
esp	4096
eip	0

don't care

cs	27
ds	35
ss	35
es	0
fs	0
gs	0

```
0  pushl $argv
    pushl $init
    pushl $0
    movl $SYS_exec,%eax
    int $64
    exit:
    movl $SYS_exit, %eax
    int $T_SYSCALL
    jmp exit

    init:
    .string "/init\0"

    .p2align 2
    argv:
    .long init
    .long 0

4095
```

The naive approach for the scheduler

```
movw $35,%ax
movw %ax,ds
movw %ax,ss # Interrupts use the stack. We just ruined
movl $4096,%esp
```

```
movl $0,%eax # Much better to move eax to other regs
movl $0,%ecx
movl $0,%edx
movl $0,%ebx
movl $0,%ebp
movl $0,%esi
movl $0,%edi
```

```
movw %ax,%es
movw %ax,%fs
movw %ax,%gs
```

```
ljmp $27,$0
```

The problem

It is not possible to find a reasonable order to load:

- cs, eip.
- ss, esp.
- eflags (Containing the Interrupt flag).

What the hardware should provide us with?

- The **iret** instruction fits the bill.

iret instruction in kernel mode

Implied stack operands:

	31		9		2	1	0
esp							
ss	unused			gdt index		0	11
eflags							
cs	unused			gdt index		0	11
esp → eip							

The operands are loaded into the appropriate registers and POPped of the stack.


Working Approach for the scheduler!!

```
pushl    $4096
pushl    $35
pushl    $512
pushl    $27
pushl    $0
```

```
iret
```


- **We are not the scheduler!**

Current stack



What do we do??

Build kernel stack for the first process

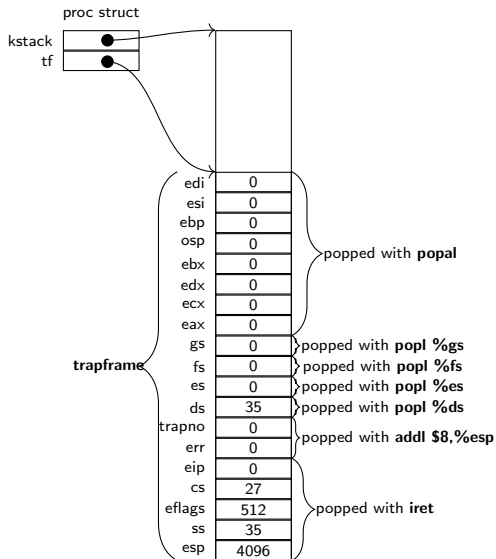


- We expect the scheduler to execute **iret** with its **esp** pointing to the **eip** field.
- We began the discussion with more registers.
- What about those other registers?
- The scheduler should load them with the prescribed values.
 - Ugly. Registers are being setup in two different places.

OR

- We will put registers values on stack for all the registers!

Build kernel stack for first process



trapframe structure

602

```
struct trapframe {  
    uint edi;  
    uint esi;  
    uint ebp;  
    uint oesp;  
    uint ebx;  
    uint edx;  
    uint ecx;  
    uint eax;  
  
    ushort gs;  
    ushort padding1;  
    ushort fs;  
    ushort padding2;
```

```
    ushort es;  
    ushort padding3;  
    ushort ds;  
    ushort padding4;  
    uint trapno;  
  
    uint err;  
    uint eip;  
    ushort cs;  
    ushort padding5;  
    uint eflags;  
  
    uint esp;  
    ushort ss;  
    ushort padding6;
```

}

Building trapframe example

```
sp      = p->kstack + KSTACKSIZE;  
sp      -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
memset(p->tf, 0, sizeof *p->tf);
```


```
p->tf->cs = (SEG_UCODE << 3) | DPL_USER;  
p->tf->ds = (SEG_UDATA << 3) | DPL_USER;  
p->tf->es = p->tf->ds;  
p->tf->ss = p->tf->ds;  
p->tf->eflags = FL_IF;  
p->tf->esp = PGSIZE;  
p->tf->eip = 0; // beginning of initcode.S
```

- So, the code the scheduler is supposed to execute is:

```
popal  
popl %gs  
popl %fs  
popl %es  
popl %ds  
addl $8,%esp  
iret
```

- For the sake of generality, we leave the address of this code on the stack.

Build kernel stack for first process



trapret pointer setting

```
sp -= 4;  
*sp = trapret;
```


A word of caution

- The **real** reason behind the **trapframe** construction is the interrupt servicing code.
- Wait for the real reasoning until there.

The **context** struct

- The **xv6** machinery is more general than the above.
- The scheduler has a context of its own.
- Using just **trapframe** will cause the scheduler to loose context.
- So, a **context** struct is pushed on the stack in order to preserve context.

Build kernel stack for the first process



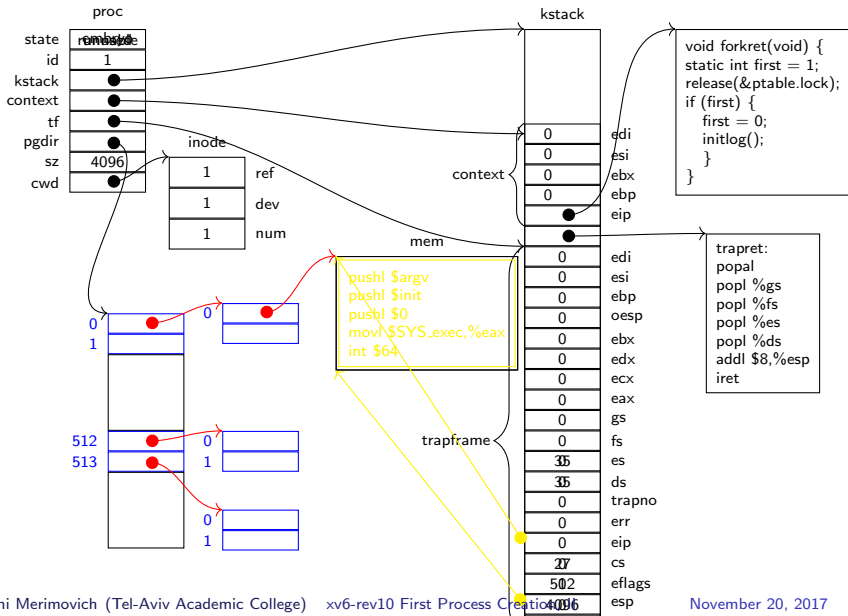
context structure

```
2623 struct context {  
    uint edi;  
    uint esi;  
    uint ebx;  
    uint ebp;  
    uint eip;  
};
```

Bulding the **context** structure

```
sp -= sizeof *p->context;  
p->context = sp;  
memset(p->context, 0, sizeof(*p->context));  
p->context->eip = forkret;
```

The whole first process creation




userinit()

- **allocproc()**: Allocate **proc** structure.
 - Allocate process kernel stack.
 - Leave space for a **trapframe** structure.
 - Build **context** structure on the stack.
- **setupkvm**: Create page table.
- **inituvm**: User space:
 - Allocate one page.
 - Copy user code to the allocated page.
 - Modify page table to use the allocated page.
- Set the **trapframe** so the user process will be able to run.

allocproc(): (1) Finding unused proc structure

```
2473 static struct proc *allocproc(void) {  
    struct proc *p;  
    char *sp;  
  
    acquire(&ptable.lock);  
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++)  
        if (p->state == UNUSED)  
            goto found;  
    release(&ptable.lock);  
    return 0;  
  
found:  
    p->state = EMBRYO;  
    p->pid = nextpid++;  
    release(&ptable.lock);
```

allocproc(): (1) Operation



```
void forkret(void) {  
    static int first = 1;  
    release(&ptable.lock);  
    if (first) {  
        first = 0;  
        initlog();  
    }  
}
```

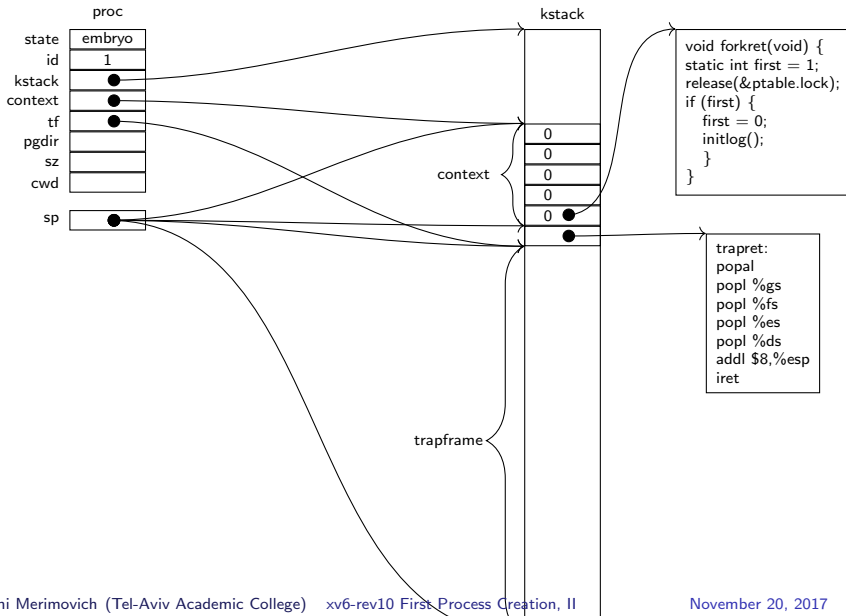
```
trapret:  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $8,%esp  
    iret
```

allocproc: (2) Initialize process kernel stack

2494

```
if ((p->kstack = kalloc()) == 0) {  
    p->state = UNUSED;  
    return 0;  
}  
sp      = p->kstack + KSTACKSIZE;  
sp      -= sizeof *p->tf;  
p->tf = (struct trapframe*)sp;  
sp      -= 4;  
*(uint*)sp = (uint)trapret;  
sp      -= sizeof *p->context;  
p->context = (struct context*)sp;  
memset(p->context, 0, sizeof *p->context);  
p->context->eip = (uint)forkret;  
return p;  
}
```

allocproc: (2) Operation




inituvm

```
1886 inituvm(pde_t *pgdir, char *init, uint sz)
{
    char *mem;

    if (sz >= PGSIZE)
        panic("inituvm: more than a page");
    mem = kalloc();
    memset(mem, 0, PGSIZE);
    mappages(pgdir, 0, PGSIZE, v2p(mem), PTE_W|PTE_U);
    memmove(mem, init, sz);
}
```


initvm() operation



userinit(): (1) Invoking allocproc & initvm

```
2520 userinit(void) {  
    struct proc *p;  
    extern char _binary_initcode_start[],  
                _binary_initcode_size[];  
  
    p = allocproc();  
    initproc = p;  
    if ((p->pgdir = setupkvm()) == 0)  
        panic("userinit: out of memory");  
    initvm(p->pgdir, _binary_initcode_start,  
            (int) _binary_initcode_size);  
    p->sz = PGSIZE;
```

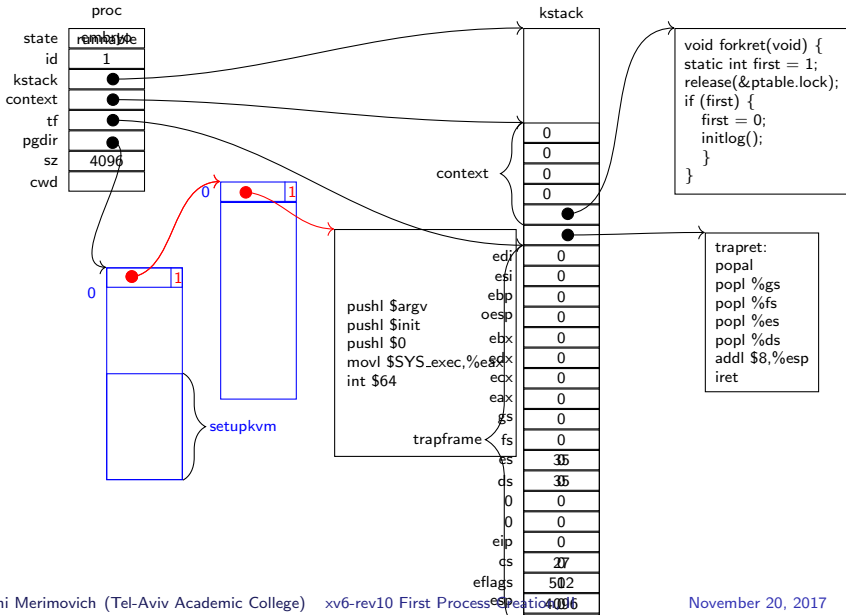

userinit(): (1) Operation



userinit(): (2) trapframe initialization

```
2532  memset(p->tf, 0, sizeof(*p->tf));
    p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
    p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
    p->tf->es = p->tf->ds;
    p->tf->ss = p->tf->ds;
    p->tf->eflags = FL_IF;
    p->tf->esp = PGSIZE;
    p->tf->eip = 0; // beginning of initcode.S
    p->cwd = namei("/");
    p->state = RUNNABLE;
}
```

userinit(): (2) Operation



trapret

```
3324 trapret :  
    popal  
    popl %gs  
    popl %fs  
    popl %es  
    popl %ds  
    addl $0x8, %esp # trapno and errcode  
    iret
```

- What is the registers state after the `iret` instruction?

Registers after iret

eax {	0	} eax,ecx,edx,ebx,ebp,esi,edi	
edi {	0		
esp {	4096		
eip {	0		
cs {	SEG_UCODE	011	
ds {	SEG_UDATA	011	
ss {	SEG_UDATA	011	
es {	0		
fs {	0		
gs {	0		

Informing the CPU about safe stack

The PUSHAL/POPAL instructions

```
pushl %eax  
pushl %ecx  
pushl %edx  
pushl %ebx  
pushl %esp !!!  
pushl %ebp  
pushl %esi  
pushl %edi
```

```
popl %edi  
popl %esi  
popl %ebp  
popl %esp  
popl %ebx  
popl %edx  
popl %ecx  
popl %eax
```

What now?

- How the kernel will return to run?
 - Only by interrupt service routine.
- Where is the interrupt service routine located?
 - Later topic,
- What happens to the stack when interrupt is delivered??
 - GOOD POINT.
 - Before going to user mode we set stack address for interrupts.

