xv6©-rev10

xv6 Interrupt Servicing

Carmi Merimovich

Tel-Aviv Academic College

December 11,2016

Interrupt Serivce routines

- In xv6 the C-routine **trap()** is the basis for interrupt servicing.
- Interrupt service routines are necessarily hardware depedent, thus
  - ISRs cannot be written **totally** in C.
  - (as usual) We aim to write whatever is possible in C.
- In the following we look at what the hardware forces us to do.

Infrastructure of interrupt service routines in C

Using **trap0()** issues

- Assume **trap0** is a C routine to handle interrupt request 0.
- Issues we need to deal with:
  - Can `IDT[0]` point directly to **trap0**?
  - What about the interrupted code registers?
  - What about our segment registers?

We solve these issues one by one.

Using **trap0()** issues

- Assume **trap0** is a C routine to handle interrupt request 0.
- Issues we need to deal with:
  - Can `IDT[0]` point directly to **trap0**? Yes, but returing will crash.
  - What about the interrupted code registers?
  - What about our segment registers?

We solve these issues one by one.

Using **trap0()** issues

- Assume **trap0** is a C routine to handle interrupt request 0.
- Issues we need to deal with:
  - Can `IDT[0]` point directly to **trap0**? Yes, but returing will crash.
  - What about the interrupted code registers? Proably will be destroyed.
  - What about our segment registers?

We solve these issues one by one.

Using **trap0()** issues

- Assume **trap0** is a C routine to handle interrupt request 0.
- Issues we need to deal with:
  - Can `IDT[0]` point directly to **trap0**? Yes, but returing will crash.
  - What about the interrupted code registers? Proably will be destroyed.
  - What about our segment registers? Unexpected if was at user mode.

We solve these issues one by one.

IDT[0] pointing to a good place

- **trap0()** ends with **ret**, no way around it.
- So we wrap **trap0** inside a routine ending with **iret**:

```
vector0:
  call trap0
  iret
```

- Then we let **IDT[0]** point to **vector0**.

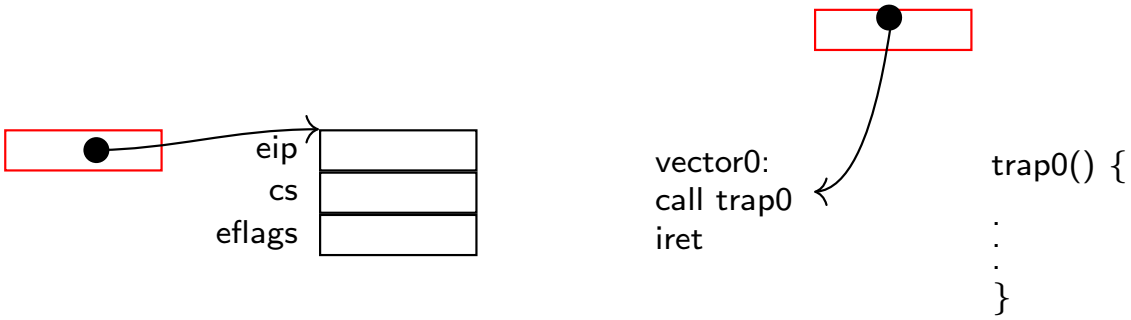
eip
cs
eflags
vector0:
call trap0
iret
trap0() {
.
.
.
}

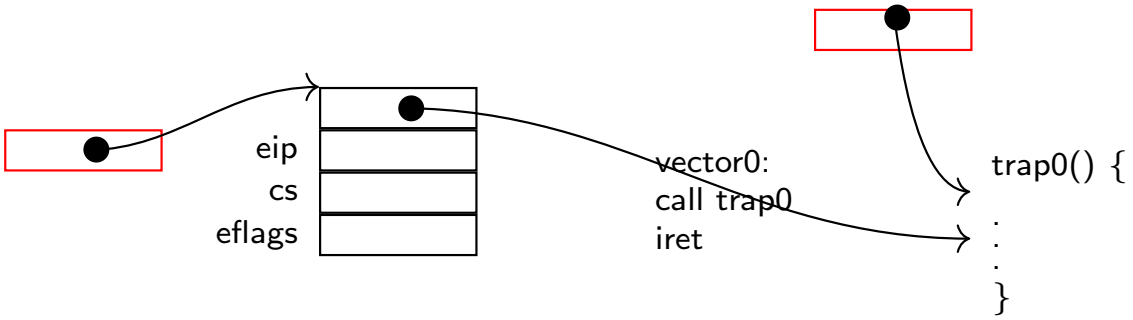
eip
cs
eflags
vector0:
call trap0
iret
trap0() {
.
.
.
}

Saving interrupted code GPRs

- Registers destroyed by the interrupt service routine should be saved.
- **eip**, **eflags**, and possibly **esp** where already saved.
- The registers destroyed by **trap0()** are not known, thus save them all:

```
vector0:
  pushal
  call trap0
  popal
  iret
```

Our segment registers?

- **cs** was saved and then loaded from the IDT.
- **ss**:
  - User mode interrupted: **ss** was saved and then loaded from the TSS.
  - Kernel mode interrupted: **ss** already contains the correct value (`SEG_KDATA`).
- **ds**, **es**, **fs**, **gs**: Code needed:

```
movw $SEG_KDATA<<3,%ax
movw %ax,%ds
movw %ax,%es
movw $0,%ax
movw %ax,%fs
movw %ax,%gs
```

Infrastructure (1) for trap0()

```
vector0:
 pushal
 movw $SEG_KDATA<<3,%ax
 movw %ax,%ds
 movw %ax,%es
 movw $0,%ax
 movw %ax,%fs
 movw %ax,%gs
```

```
 call trap0
 popal
 iret
```

- We might have just destroyed user mode registers!

Infrastructure (2) for trap0()

```
vector0:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw  $SEG_KDATA<<3,%ax
 movw  %ax,%ds
 movw  %ax,%es
 movw  $0,%ax
 movw  %ax,%fs
 movw  %ax,%gs
```

```
 call  trap0
 popal
 popl %gs
 popl %fs
 popl %es
 popl %ds
 iret
```

Interrupted code registers content

- ISRs might be interested in accessing the interrupted code registers.
- (This is reasonable only for internal interrupt requests!)
- The register values where pushed onto the stack.
- A slight addition allows easy access to them from C.
- (elective) The values are accessible also without the change.

Infrastructure (3) for trap0()

```
vector0:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw  $SEG_KDATA<<3,%ax
 movw  %ax,%ds
 movw  %ax,%es
 movw  $0,%ax
 movw  %ax,%fs
 movw  %ax,%gs
```

```
pushl %esp
call  trap0
addl  $4,%esp
popal
popl %gs
popl %fs
popl %es
popl %ds
iret
```

frame struct

```
struct frame {
 long edi;
 long esi;
 long ebp;
 long esp_irrelevant;
 long ebx;
 long edx;
 long ecx;
 long eax;

 short gs;
 short padding1;
 short fs;
 short padding2;
 short es;
 short padding3;
 short ds;
 short padding4;

 long  eip;
 short cs;
 short padding5;
 long eflags;

 long esp;
 short ss;
 short padding6;
}
```

Accessing interrupted code registers

- Possible declaration of `trap0()`:

```
void trap0(struct frame *f) {
 f->eax = someValue(f);
 ⋮
}
```

- Note: The **esp** and **ss** fields exist only if user mode was interrupted.

Infrastructure for **trap8()** (serving IRQ 8)

- IRQ8 leaves error code below **eip**.
- Hence:
  - The **frame** struct does not give the correct picture.
  - There is a superfulous word before the **iret** instruction.
- So:
  - A new structure replacing **frame** struct is needed.
  - Revmoing a word from the stack is needed.

Infrastructure (1) for **trap8()**

```
vector8:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw  $SEG_KDATA<<3,%ax
 movw  %ax,%ds
 movw  %ax,%es
 movw  $0,%ax
 movw  %ax,%fs
 movw  %ax,%gs
```

```
 pushl  %esp
 call  trap8
 addl  $4,%esp
 popal
 popl %gs
 popl %fs
 popl %es
 popl %ds
 addl  $4,%esp
 iret
```

eframe struct

```
struct eframe {
 long edi;
 long esi;
 long ebp;
 long esp_irrelevant;
 long ebx;
 long edx;
 long ecx;
 long eax;

 short gs;
 short padding1;
 short fs;
 short padding2;
 short es;
 short padding3;
 short ds;
 short padding4;

 long err;
 long eip;
 short cs;
 short padding5;
 long eflags;

 long esp;
 short ss;
 short padding6;
}
```

Accessing error code

- Possible declaration of `trap8()`:

```
void trap8(struct eframe *f) {
 if (f->err)
  f->eax = someValue(f);
 else
  f->eax = someOtherValue(f);
 ⋮
}
```

Uniform behaviour would be useful

- The above solution is feasible.
- However, if one uniform routine `trap()` is asked for, we tinker.
- How?

Infrastructure (4) for **trap0()**

```
vector0:
 pushl  $0
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw  $SEG_KDATA<<3,%ax
 movw  %ax,%ds
 movw  %ax,%es
 movw  $0,%ax
 movw  %ax,%fs
 movw  %ax,%gs
```

```
 pushl  %esp
 call  trap0
 addl  $4,%esp
 popal
 popl  %gs
 popl  %fs
 popl  %es
 popl  %ds
 addl  $4,%esp
 iret
```

The **trapframe** struct (1)

```
struct trapframe {
 long edi;
 long esi;
 long ebp;
 long esp_irrelevant;
 long ebx;
 long edx;
 long ecx;
 long eax;

 short gs;
 short padding1;
 short fs;
 short padding2;
 short es;
 short padding3;
 short ds;
 short padding4;

 long  err;
 long  eip;
 short cs;
 short padding5;
 long  eflags;

 long  esp;
 short ss;
 short padding6;
}
```

All isr's can use the **trapframe** struct

```
void trap0(struct trapframe *f) {
  ⋮
}

void trap8(struct trapframe *f) {
  ⋮
}
```

- Let us reorganize the entry code.

Infrastructure (5) for **trap0()**

```
vector0:
 pushl  $0
 jmp alltraps

alltraps:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw $SEG_KDATA<<3,%ax
 movw %ax,%ds
 movw %ax,%es
 movw $0,%ax
 movw %ax,%fs
 movw %ax,%gs
```

```
pushl  %esp
call trap0
addl $4,%esp
popal
popl %gs
popl %fs
popl %es
popl %ds
addl $4,%esp
iret
```

Infrastructure (5) for **trap8()**

```
vector8:
 jmp alltraps


alltraps:
 pushl %ds
 pushl %es
 pushl %fs
 pushl %gs
 pushal
 movw $SEG_KDATA<<3,%ax
 movw %ax,%ds
 movw %ax,%es
 movw $0,%ax
 movw %ax,%fs
 movw %ax,%gs
```

```
pushl %esp
call trap8
addl $4,%esp
popal
popl %gs
popl %fs
popl %es
popl %ds
addl $4,%esp
iret
```

Infrastructure (1) for **trap()**, IRQ0

```
vector0:
 pushl  $0
 jmp alltraps
vector8:
 jmp alltraps
alltraps:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw $SEG_KDATA<<3,%ax
```

```
movw %ax,%ds
movw %ax,%es
movw $0,%ax
movw %ax,%fs
movw %ax,%gs
pushl  %esp
call trap
addl $4,%esp
popal
popl %gs
popl %fs
popl %es
popl %ds
addl $4,%esp
iret
```

Infrastructure (1) for **trap()**, IRQ0

```
vector0:
 pushl  $0
 jmp alltraps
vector8:
 jmp alltraps
alltraps:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw $SEG_KDATA<<3,%ax
```

What is lost??

```
movw %ax,%ds
movw %ax,%es
movw $0,%ax
movw %ax,%fs
movw %ax,%gs
pushl  %esp
call trap
addl $4,%esp
popal
popl %gs
popl %fs
popl %es
popl %ds
addl $4,%esp
iret
```

Infrastructure (2) for **trap()**

```
vector0:
 pushl  $0
 pushl  $0
 jmp alltraps
vector8:
 pushl  $8
 jmp alltraps
alltraps:
 pushl  %ds
 pushl  %es
 pushl  %fs
 pushl  %gs
 pushal
 movw $SEG_KDATA<<3,%ax
```

```
movw %ax,%ds
movw %ax,%es
movw $0,%ax
movw %ax,%fs
movw %ax,%gs
pushl  %esp
call trap
addl $4,%esp
popal
popl %gs
popl %fs
popl %es
popl %ds
addl $8,%esp
iret
```

The **trapframe** struct (2)

```
struct trapframe {
 long edi;
 long esi;
 long ebp;
 long esp_irrelevant;
 long ebx;
 long edx;
 long ecx;
 long eax;

 short gs;
 short padding1;
 short fs;
 short padding2;
 short es;
 short padding3;
 short ds;
 short padding4;

 long trapno;
 long err;

 long eip;
 short cs;
 short padding5;
 long eflags;

 long esp;
 short ss;
 short padding6;
}
```

trap()

```
void trap(struct trapframe *tf) {
 switch (tf->trapno) {
 case 0:
  trap0(tf);
  break;

  ⋮

 case 8:
  trap8(tf);
  break;

 ⋮
 }
}
```

Retro groking on process creation

- Note the **trapframe** structure and recall **userinit()**.
- Now we can understand why things are as they are:
  - There are always interrupts.
  - User mode code is always going to be interrupted.
  - So there must be a framework for:
    - switching from user mode to kernel mode.
    - Return safely from kernel mode to user mode.
  - Running new user code builds on this framework.
  - New user code has no history (It is new!).
    - So a fake history is build for it.
    - And then the kenel 'returns' to the code entry point.

- The xv6 code.

Kernel entry points

We combine the previous techniques to get:

- Kernel entry points **vector0** upto **vector255** follows on the next slide.
- Each entry will contains one or two **push** instructions followed by

```
jmp  alltraps
```

- Use some scripting language to generate the `.S` file with the above code.
- (elective) (or by a smart.s and some fancy assembly directives).

all vectors

```
vector0:
 pushl  $0
 pushl  $0
 jmp  alltraps
  ⋮
vector7:
 pushl  $0
 pushl  $7
 jmp  alltraps
vector8:
 pushl  $8
 jmp  alltraps
vector9:
 pushl  $0
 pushl  $9
 jmp  alltraps
```

```
vector10:
 pushl  $10
 jmp  alltraps
 ⋮
vector14:
 pushl  $14
 jmp  alltraps
vector15:
 pushl  $0
 pushl  $15
 jmp  alltraps
vector16:
 pushl  $0
 pushl  $16
 jmp  alltraps
```

```
vector17:
 pushl  $17
 jmp  alltraps
vector18:
 pushl  $0
 pushl  $18
 jmp  alltraps
  ⋮
vector255:
 pushl  $0
 pushl  $255
 jmp  alltraps
```

all vectors

```
vector0:
 pushl   $0
 pushl   $0
 jmp   alltraps
  ⋮
vector7:
 pushl   $0
 pushl   $7
 jmp   alltraps
vector8:
 pushl   $8
 jmp   alltraps
vector9:
 pushl   $0
 pushl   $9
 jmp   alltraps
```

```
vector10:
 pushl   $10
 jmp   alltraps
 ⋮
vector14:
 pushl   $14
 jmp   alltraps
vector15:
 pushl   $0
 pushl   $15
 jmp   alltraps
vector16:
 pushl   $0
 pushl   $16
 jmp   alltraps
```

```
vector17:
 pushl   $17
 jmp   alltraps
vector18:
 pushl   $0
 pushl   $18
 jmp   alltraps
  ⋮
vector255:
 pushl   $0
 pushl   $255
 jmp   alltraps
```

Tedious!

alltraps

```
alltraps::
  pushl  %ds
  pushl  %es
  pushl  %fs
  pushl  %gs
  pushal
  movw  $SEG_KDATA<<3,%ax
  movw  %ax,%ds
  movw  %ax,%es
  movw  $0,%ax
  movw  %ax,%fs
  movw  %ax,%gs
  pushl %esp
  call  trap
  addl  $4,%esp
trapret:
  popal
  popl %gs
  popl %fs
  popl %es
  popl %ds
  addl  $8,%esp
  iret
```

vectors[]

- How do we loop on **vector0** through **vector255** to initialize IDT?

vectors[]

- How do we loop on **vector0** through **vector255** to initialize IDT?
  - We cannot.

vectors[]

- How do we loop on **vector0** through **vector255** to initialize IDT?
  - We cannot.
- So, we build a vector of addresses:

```
vectors::
  .long vector0
  .long vector1
  ⋮
  .long vector255
```

- Writing this is also TEDIOUS.

vectors.pl

```
for (my $i = 0; $i < 256; $i++) {
 print "vector$i:\n";
 if (!($i == 8 ||
     ($i >= 10 && $i <= 14) || $i == 17)) {
  print " pushl \$0\n";
 }
 print " pushl \$$i\n";
 print " jmp alltraps\n";
}
print ".data\n";
print "vectors::\n";
for (my $i = 0; $i < 256; $i++) {
 print " .long vector$i\n";
}
```

(elective) a direct vectors.S

```
 .globl  alltraps, vectors
 .data
vectors:
 v = 0
.rept 256
 .text
0:
 .if !(v == 8 || (10 <= v && v <= 14) || v == 17)
   pushl $0
 .endif
 pushl $v
 jmp alltraps
 .data
 .long 0b
 v = v + 1
.endr
.end
```

IDT initializing

Gate descriptor

Trap/Interrupt gate:

31 … 16	15 (P)	14 13 (DPL)	12 (S)	11 … 8 (Type)	7 … 0
Offset (high)	1	00	0	111 1/0	00000000

31 … 16	15 … 0
Selector	Offset (low)

```
struct gatedesc {
  uint off_15_0 : 16; // low 16 bits of offset in segment
  uint cs : 16; // code segment selector
  uint args : 5; // # args, 0 for interrupt/trap gates
  uint rsv1 : 3; // reserved(should be zero I guess)
  uint type : 4; // type(STS_{TG,IG32,TG32})
  uint s : 1; // must be 0 (system)
  uint dpl : 2; // descriptor(meaning new) privilege level
  uint p : 1; // Present
  uint off_31_16 : 16; // high bits of offset in segment
};
```

SETGATE macro

```
921 #define SETGATE(gate, istrap, sel, off, d) \ { \
      (gate).off_15_0 = (uint)(off) & 0xffff; \
      (gate).cs = (sel); \
      (gate).args = 0; \
      (gate).rsv1 = 0; \
      ((gate).type = (istrap) ? STS_TG32 : STS_IG32; \
      (gate).s = 0; \
      (gate).dpl = (d); \
      (gate).p = 1; \
      (gate).off_31_16 = (uint)(off) >> 16; \
    }
```

IDT construction

```
#define T_SYSCALL 64

3361 struct gatedesc idt[256];
     extern uint vectors[]; // in vectors.S: array of 256

3367 tvinit(void) {
      for (int i = 0; i < 256; i++)
       SETGATE(idt[i], 0, SEG_KCODE<<3, vectors[i], 0);

      SETGATE(idt[T_SYSCALL], 1, SEG_KCODE<<3,
                  vectors[T_SYSCALL], DPL_USER);
     }

3378 void idtinit(void) {
      lidt(idt, sizeof(idt));
     }
```

Finaly, we can investigate **trap()**.

trap()

The code of **trap()** is composed of three parts:

1. Handle **killed** and dispatch to **syscall()** if invoked by **int $64**.
2. Dispatch to code handling other IRQ's.
3. Code handling **killed** and **yield()**.

Note about **killed**:

- **killed** set means the process should terminate.
- The **killed** field should be checked at the following locations:
  - Entering the kernel.
  - Leaving the kernel.
  - Whenever the process might enter **SLEEPING** for a long time.

trap() part 1

```
3401 void trap(struct trapframe *tf) {
      if (tf->trapno == T_SYSCALL) {
        if (myproc()->killed)
          exit();
        myproc()->tf = tf;
        syscall();
        if (myproc()->killed)
          exit();
        return;
      }
```

- **syscall()** might be long (i.e., invokes **swtch**), hence the double check on **myproc()->killed**.

trap() part 2, controller interrupts

```
3413  switch (tf->trapno) {
      case T_IRQ0+IRQ_TIMER:
       if (cpuid() == 0) {
        acquire(&tickslock);
        ticks++;
        wakeup(&ticks);
        release(&tickslock);
       }
       lapiceoi();
       break;
      case T_IRQ0+IRQ_IDE:
       ideintr();
       lapiceoi();
       break;
      case T_IRQ0+IRQ_IDE+1:
       break;
```

```
      case T_IRQ0+IRQ_KBD:
       kbdintr();
       lapiceoi();
       break;
      case T_IRQ0+IRQ_COM1:
       uartintr();
       lapiceoi();
       break;
      case T_IRQ0+7:
      case T_IRQ0+IRQ_SPURIOUS:
       cprintf("cpu%d: spurious \
    interrupt at %x:%x\n",
        cpuid(), tf->cs, tf->eip);
       lapiceoi();
       break;
```

trap() part 2, unexpected interrupt

```
3450    default:
         if (myproc() == 0 || (tf->cs&3) == 0) {
          cprintf("unexpected trap %d from cpu %d \
        eip %x (cr2=0x%x)\n",
                  tf->trapno, mycpu()->id, tf->eip, rcr2());
          panic("trap");
         }
         cprintf("pid %d %s: trap %d err %d on cpu %d "
              "eip 0x%x addr 0x%x      kill proc\n",
         myproc()->pid, myproc()->name, tf->trapno, tf->err
                  cpuid(), tf->eip,
         rcr2());
         myproc()->killed = 1;
        }
```

trap() part 3

```
3468    if (myproc() && myproc()->killed &&
                (tf->cs&3) == DPL_USER)
         exit();

        if (myproc() && myproc()->state == RUNNING &&
                tf->trapno == T_IRQ0+IRQ_TIMER)
         yield();

        if (myproc() && myproc()->killed &&
                (tf->cs&3) == DPL_USER)
         exit();
    }
```