

xv6©-rev10  
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)  
sleep syscall

Carmi Merimovich

Tel-Aviv Academic College

February 20, 2017

- The `sleep` system call suspends execution of the process for the number of ticks supplied by the argument.
- There is an argument!!
- It must be checked carefully!!
- The `sys_sleep` implementation is very simple:
  - It is assumed each clock tick declares an event with id `&ticks`.
  - `sys_sleep` waits for the `&ticks` event.
  - When `sys_sleep` resumes execution, it checks if it was suspended for long enough.
  - If not it returns to the event waiting.

## Variables in `sys_sleep`

- `n`: Number of ticks to wait.
- `ticks`: Global variable containing the number of ticks from boot.
- `tickslock`: A spinlock protecting `ticks`.

## sys\_sleep

3815

```
sys_sleep(void) {  
    int n;  
    uint ticks0;  
  
    if (argint(0, &n) < 0) return -1;  
    acquire(&tickslock);  
    ticks0 = ticks;  
    while (ticks - ticks0 < n) {  
        if (myproc() -> killed) {  
            release(&tickslock);  
            return -1;  
        }  
        sleep(&ticks, &tickslock);  
    }  
    release(&tickslock);  
    return 0;  
}
```

## ticks?!

Somewhere in the code the following should happen:

- ticks increments.
- Event &ticks is declared.

In trap():

```
3414  case T_IRQ0+IRQ_TIMER:
      if (cpuid() == 0) {
        acquire(&tickslock);
        ticks++;
        wakeup(&ticks);
        release(&tickslock);
      }
      lapiceoi();
      break;
```