

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
First Process Creation, I

Carmi Merimovich

Tel-Aviv Academic College

November 10, 2017

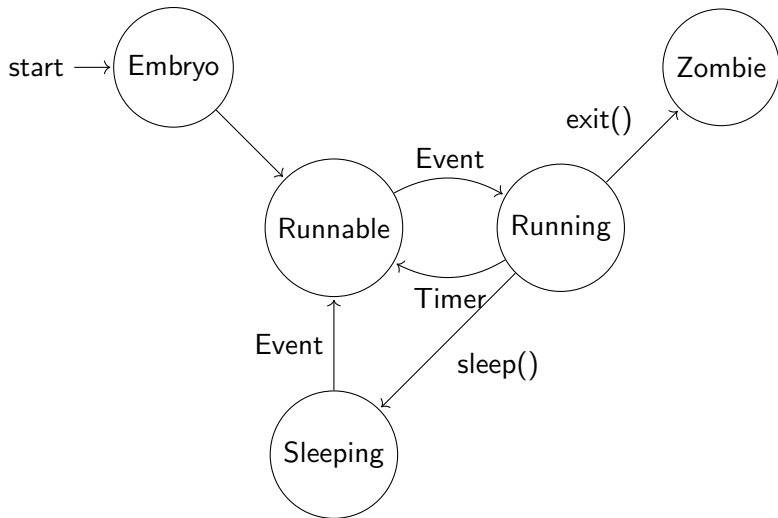
Context

```
kinit1(end, P2V(4*1024*1024)); // phys page allocation
kvmalloc(); // kernel page table
:
segininit(); // set up segments
:
pinit(); // process table
:
:
kinit2(P2V(4*1024*1024), P2V(PHYSTOP)); // must copy
userinit(); // first user process
mpmain();
```

struct proc

```
2337 struct proc {  
    uint sz; // Size of process memory (bytes)  
    pde_t* pgdir; // Page table  
    char *kstack; // Bottom of kernel stack for this pr  
    enum procstate state; // Process state  
    volatile int pid; // Process ID  
    struct proc *parent; // Parent process  
    struct trapframe *tf; // Trap frame for current sys  
    struct context *context; // swtch() here to run pro  
    void *chan; // If non zero, sleeping on chan  
    int killed; // If non zero, have been killed  
    struct file *ofile[NOFILE]; // Open files  
    struct inode *cwd; // Current directory  
    char name[16]; // Process name (debugging)  
};
```

state transition



All

transition are due to INTERRUPTS.

proc structure storage

```
2334 enum procstate { UNUSED, EMBRYO, SLEEPING,  
                                RUNNABLE, RUNNING, ZOMBIE }  
  
2409 struct {  
    struct spinlock lock;  
    struct proc proc[NPROC];  
} ptable;
```

First process simplifying assumptions

- Only 4KB user space area.
- The code/data content is stored at some static place in the kernel.
- User code begins running at address 0.
- Top of user stack shall be at 4096.

Initial user mode state for first process

eax	0
ebx	0
ecx	0
edx	0
ebp	0
esi	0
edi	0
esp	4096
eip	0

don't care

cs	27
ds	35
ss	35
es	0
fs	0
gs	0

```
0  pushl $argv
    pushl $init
    pushl $0
    movl $SYS_exec,%eax
    int $64
    exit:
    movl $SYS_exit, %eax
    int $T_SYSCALL
    jmp exit

    init:
    .string "/init\0"

    .p2align 2
    argv:
    .long init
    .long 0

4095
```

xv6 steps taken to construct the state above

1. Allocating memory for user code and setting its content.
2. Adding mapping rules.
3. Store on the kernel stack the expected user mode registers values.

Putting initial user code into kernel data vector

initcode.S: User mode code of 1st process

8408

```
#include "syscall.h"
#include "traps.h"

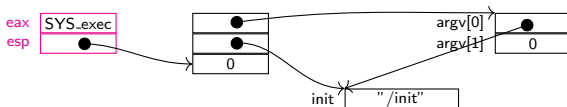
.globl start
start:
    pushl $argv
    pushl $init
    pushl $0 // caller pc
    movl  $SYS_exec,%eax
    int   $T_SYSCALL

exit:
    movl  $SYS_exit, %eax
    int  $T_SYSCALL
    jmp  exit
```

```
init:
    .string "/init\0"

    .p2align 2
argv:
    .long init
    .long 0
```

State on the first **int \$64**



C code equivalent

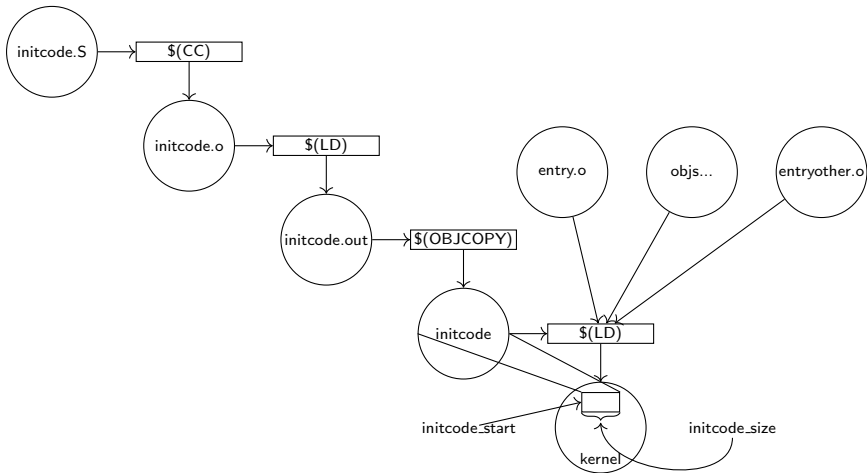
```
char *argv[] = {"/init", 0};  
main() {  
    exec(argv[0], argv);  
    for (;;)   
        exit{};  
}
```

initcode.S: Compiling, Linking, Embedding

```
initcode: initcode.S
$(CC) $(CFLAGS) -nostdinc -I. -c initcode.S
$(LD) $(LDFLAGS) -N -e start -Ttext 0 \
        -o initcode.out initcode.o
$(OBJCOPY) -S -O binary initcode.out initcode

kernel: $(OBJS) entry.o entryother initcode kernel.ld
$(LD) $(LDFLAGS) -T kernel.ld -o kernel \
        entry.o $(OBJS) -b binary initcode entryother
```

Diagram of compiling, linking, embedding

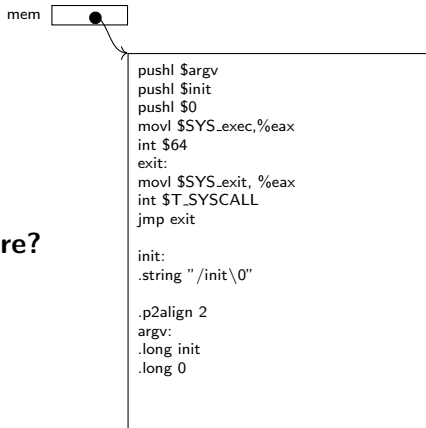


Building the memory area for the user mode.

Building user memory

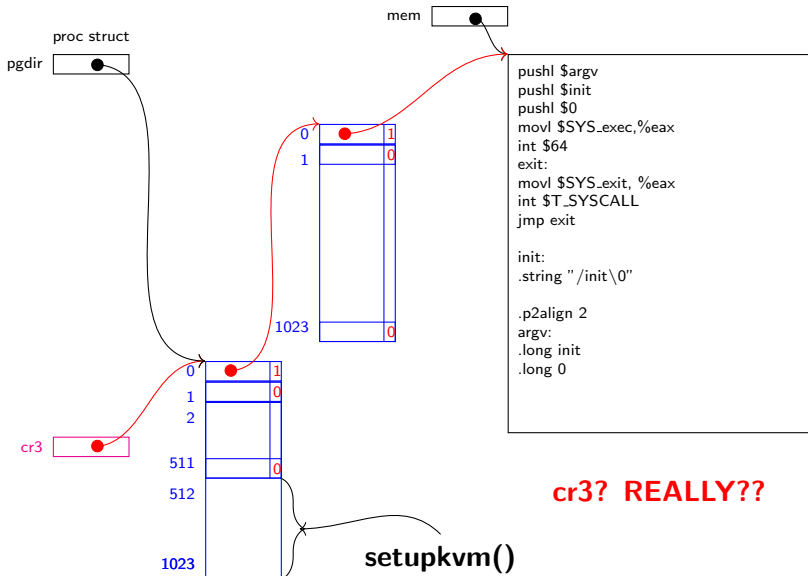
```
extern char initcode_start[], initcode_size[];  
mem = kalloc();  
memset(mem, 0, PGSIZE);  
memmove(mem, initcode_start, initcode_size);
```


Building user memory

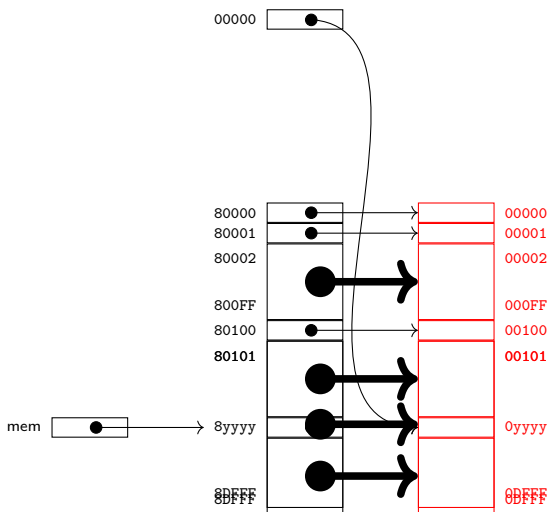


Do we have user mode addresses here?

Mapping the first 4KB of user space



User mode pages are doubly mapped



Mapping first user page!

```
int *pgdir = setupkvm();
```

```
int *pte = walkpgdir(pgdir, 0, 1);  
*pte = v2p(mem) | (PTE_P|PTE_U|PTE_W);
```