

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
System calls

Carmi Merimovich

Tel-Aviv Academic College

November 30, 2017

syscall()

- **syscall()** dispatches to the function implementing the relevant system call.
- System call number in user mode is in **eax**.
- Hence it is in **myproc()->tf->eax**.
- We will be done by having **syscalls[]**, a vector of routines addresses.
- return value of **syscalls[myproc()->tf->eax]** is put into **myproc()->tf->eax**.

syscall.h

3201

```
#define SYS_fork 1
#define SYS_exit 2
#define SYS_wait 3
#define SYS_pipe 4
#define SYS_read 5
#define SYS_kill 6
#define SYS_exec 7
#define SYS_fstat 8
#define SYS_chdir 9
#define SYS_dup 10
#define SYS_getpid 11
#define SYS_sbrk 12
#define SYS_sleep 13
#define SYS_uptime 14
#define SYS_open 15
#define SYS_write 16
```

```
#define SYS_mknod 17
#define SYS_unlink 18
#define SYS_link 19
#define SYS_mkdir 20
#define SYS_close 21
```

syscalls[]

3350

```
static int (*syscalls[])(void) = {  
    [SYS_fork]    sys_fork ,  
    [SYS_exit]    sys_exit ,  
    [SYS_wait]    sys_wait ,  
    [SYS_pipe]    sys_pipe ,  
    [SYS_read]    sys_read ,  
    [SYS_kill]    sys_kill ,  
    [SYS_exec]    sys_exec ,  
    [SYS_fstat]   sys_fstat ,  
    [SYS_chdir]   sys_chdir ,  
    [SYS_dup]     sys_dup ,  
    [SYS_getpid]  sys_getpid ,  
    [SYS_sbrk]    sys_sbrk ,  
    [SYS_sleep]   sys_sleep ,  
    [SYS_uptime]  sys_uptime ,  
    [SYS_open]    sys_open ,  
    [SYS_write]   sys_write ,  
    [SYS_mknod]   sys_mknod ,  
    [SYS_unlink]  sys_unlink ,  
    [SYS_link]    sys_link ,  
    [SYS_mkdir]   sys_mkdir ,  
    [SYS_close]   sys_close ,  
};
```

syscall()

```
3701 void syscall(void) {  
    int num;  
  
    num = myproc()->tf->eax;  
    if (num > 0 && num < NELEM(syscalls) &&  
        syscalls[num]) {  
        myproc()->tf->eax = syscalls[num]();  
    } else {  
        cprintf("%d_%s: unknown_sys_call_%d\n",  
            myproc()->pid, myproc()->name, num);  
        myproc()->tf->eax = -1;  
    }  
}
```

myproc()->tf->eax = syscalls[num]()

```
switch (num) {  
case SYS_fork:  
    myproc()->tf->eax = sys_fork();  
    break;  
case SYS_exit:  
    myproc()->tf->eax = sys_exit();  
    break;  
case SYS_mkdir:  
    myproc()->tf->eax = sys_mkdir();  
    break;  
case SYS_close:  
    myproc()->tf->eax = sys_close();  
    break;  
default:  
    myproc()->tf->eax = -1;  
    break
```

We can investigate each system call now...

Process control syscalls

Our current knowledge is enough for the first two syscalls:

- **getpid()**.
- **fork()**.

In addition, for the following syscalls we need to access arguments:

- **kill()**.

In addition, for the following we also need the event system:

- **exit()**.
- **wait()**.
- **sleep()**.

In addition, for the following syscall we need to read file:

- **exec()**.