

xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
exec syscall II

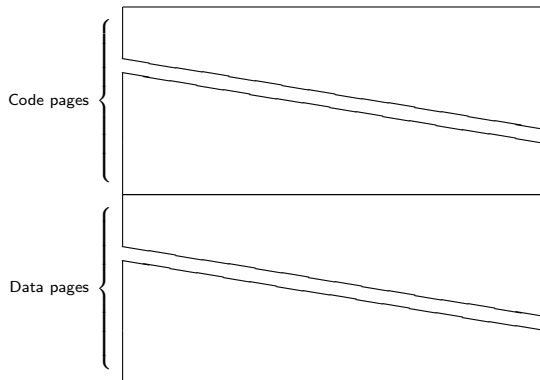
Carmi Merimovich

Tel-Aviv Academic College

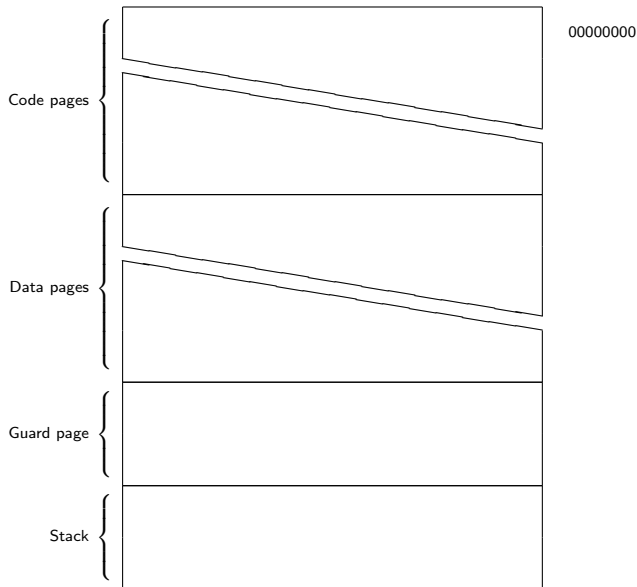
January 10, 2017

exec: step 2.

Address space after loading elf



Adding guard page and stack page



exec: Allocate user stack and guard page

6664

```
sz = PGROUNDUP(sz);  
if ((sz = allocuvm(pgdir, sz, sz + 2*PGSIZE)) == 0)  
    goto bad;  
clearpteu(pgdir, (char*)(sz - 2*PGSIZE));  
sp = sz;
```

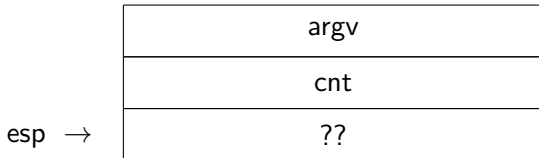
- Note the sp variable for later use.

exec: step 3

- The invoked executable expects to start with the following signature:

`main(int cnt , char **argv);`

- Hence the following stack should be prepared:



- Hence somewhere a vector of pointers should be allocated.
- Hence a copy of the strings supplied to the exec syscall should be done.
- All these allocation are done on the new stack.
- Note, **I believe** all this should be done in user mode!

Copying between address spaces: Common AS

- First tool: Convert full virtual address to kernel address.
- Note 0 is not a legal kernel address.

```
char *my_uva2ka(pde_t *pgdir, char *uva) {  
    pte_t *pte = walkpgdir(pgdir, uva, 0);  
    if (pte == 0 || (*pte & PTE_P) == 0)  
        return 0;  
    if ((*pte & PTE_U) == 0)  
        return 0;  
    return (char*)P2V(PTE_ADDR(*pte) |  
                      (uva & (PGSIZE - 1)));  
}
```


Copying between address spaces: Generalized addresses

```
int gcopy(pde_t *tpgdir, int taddr,
          pde_t *fpgdir, int faddr, int len) {
    for (int i = 0; i < len; i++) {
        char *fka = my_uva2ka(fpgdir, faddr++);
        char *tka = my_uva2ka(tpgdir, taddr++);
        if (fka == 0 || tka == 0) return 0;
        *tka = *fka;
    }
}
```

- The above code is horribly inefficient.
- However, it is quite clear.

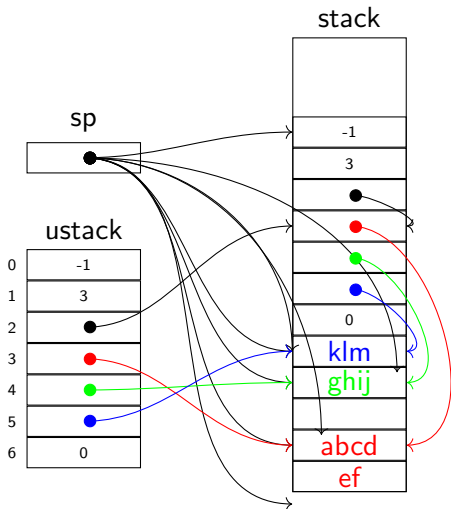
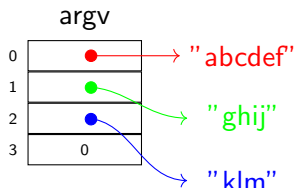
xv6 copying between address spaces: Common AS

- Convert full page address to kernel address.
- Note 0 is not a legal kernel address.

```
2102 char *uva2ka(pde_t *pgdir, char *uva) {  
    pte_t *pte = walkpgdir(pgdir, uva, 0);  
    if ((*pte & PTE_P) == 0)  
        return 0;  
    if ((*pte & PTE_U) == 0)  
        return 0;  
    return (char*)P2V(PTE_ADDR(*pte));  
}
```

xv6 copy to a generalized address

```
2118 int copyout(pde_t *pgdir, uint va,
              void *p, uint len) {
    char *buf = (char*)p;
    while (len > 0) {
        uint va0 = (uint)PGROUNDDOWN(va);
        char *pa0 = uva2ka(pgdir, (char*)va0);
        if (pa0 == 0) return -1;
        uint n = PGSIZE - (va - va0);
        if (n > len) n = len;
        memmove(pa0 + (va - va0), buf, n);
        len -= n;
        buf += n;
        va = va0 + PGSIZE;
    }
    return 0;
}
```



exec: Copying argv[]

6671

```
uint  ustack[3+MAXARG+1];
for (argc = 0; argv[argc]; argc++) {
    if (argc >= MAXARG) goto bad;
    sp = (sp - (strlen(argv[argc]) + 1)) & ~3;
    if (copyout(pgdir, sp, argv[argc],
                strlen(argv[argc]) + 1) < 0) goto bad;
    ustack[3+argc] = sp;
}
ustack[3+argc] = 0;

ustack[0] = 0xffffffff; // fake return PC
ustack[1] = argc;
ustack[2] = sp - (argc+1)*4; // argv pointer

sp -= (3+argc+1) * 4;
if (copyout(pgdir, sp, ustack, (3+argc+1)*4) < 0)
    goto bad;
```

exec: step 4

exec: Switching address space, fixing the **trapframe**

6690

```
for (last=s=path; *s; s++) if (*s == '/')
    last = s+1;
safestrcpy(myproc()->name, last, sizeof(myproc()->n

oldpgdir = myproc()->pgdir;
myproc()->pgdir = pgdir;
myproc()->sz = sz;
myproc()->tf->eip = elf.entry; // main
myproc()->tf->esp = sp;
switchvm(myproc());
freevm(oldpgdir);
return 0;
bad:
if (pgdir) freevm(pgdir);
if (ip) {iunlockput(ip);end_op();}
return -1;
```