

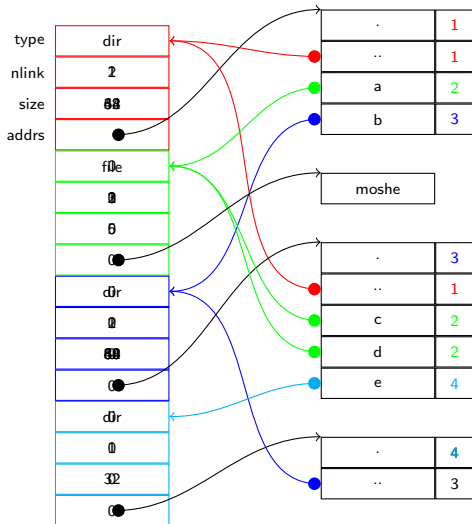
xv6©-rev10
(Copyright Frans Kaashoek, Robert Morris, and Russ Cox.)
Name Layer I

Carmi Merimovich

Tel-Aviv Academic College

January 20, 2017

dinodes and directories on disk



opendir("/usr/bin/ls")

System calls using the name layer

- `sys_open.`
- `sys_mkdir.`
- `sys_link.`
- `sys_unlink.`

open(char *path, int mode)

- omode flags:

```
3950 #define O_RDONLY 0x000
      #define O_WRONLY 0x001
      #define O_RDWR 0x002
      #define O_CREATE 0x200
```

- If O_CREATE set then we need to create file.
- If O_CREATE is clear we need to open existing file.
- Folders can be opened only for readonly access.

sys_open logic

- If creating is attempted then:
 - Delegate to the create function in the inode layer.
 - (Thus, either creating a file, or opening an existing file).
 - Note create an ilock'ed inode pointer.
- If opening is attempted then:
 - Delegate to the namei function in the inode layer.
 - (Thus opening an existing file).
 - If the file is directory make sure readonly access is in effect.
- Build a FD_INODE file structure pointing to the found/created inode.
- Hide the file pointer in the ofile vector.

sys_open, invoking lower layer

```
6401 int sys_open(void) {  
    char *path;  
    int omode;  
    struct inode *ip;  
    if (argstr(0, &path) < 0 || argint(1, &omode) < 0) return  
        begin_op();  
    if (omode & O_CREATE) {  
        if ((ip = create(path, T_FILE, 0, 0)) == 0) {  
            end_op(); return -1;  
        }  
    } else {  
        if ((ip = namei(path)) == 0) { end_op(); return -1;}  
        ilock(ip);  
        if (ip->type == T_DIR && omode != O_RDONLY) {  
            iunlockput(ip);  
            end_op();  
            return -1;  
        }  
    }  
}
```

sys_open, housekeeping

6432

```
if ((f=filealloc()) == 0 || (fd=fdalloc(f)) < 0) {  
    if (f)  
        fileclose(f);  
    iunlockput(ip);  
    end_op()  
    return -1;  
}  
iunlock(ip);  
end_op();  
  
f->type = FD_INODE;  
f->ip = ip;  
f->off = 0;  
f->readable = !(omode & O_WRONLY);  
f->writable = (omode & O_WRONLY) || (omode & O_RDWR);  
return fd;
```

sys_mkdir

```
6451 int sys_mkdir(void) {  
    char *path;  
    struct inode *ip;  
  
    begin_op();  
    if (argstr(0, &path) < 0 ||  
        (ip = create(path, T_DIR, 0, 0)) == 0) {  
        end_op();  
        return 1;  
    }  
    iunlockput(ip);  
    end_op();  
    return 0;  
}
```


sys_link (1)

```
6202 sys_link(void) {  
    char name[DIRSIZ], *new, *old;  
    struct inode *dp, *ip;  
  
    if (argstr(0, &old) < 0 || argstr(1, &new) < 0)  
        return -1;  
    begin_op()  
    if ((ip = namei(old)) == 0) {  
        end_op();  
        return -1;  
    }  
    ilock(ip);  
    if (ip->type == T_DIR) {  
        iunlockput(ip);  
        end_op();  
        return -1;  
    }
```

sys_link (2)

```
6223 ip->nlink++;
    iupdate(ip);
    iunlock(ip);

    if ((dp = nameiparent(new, name)) == 0)
        goto bad;
    ilock(dp);
    if (dp->dev != ip->dev || dirlink(dp, name, ip->inum)
        iunlockput(dp);
        goto bad;
    }
    iunlockput(dp);
    iput(ip);

end_op();
return 0;
```

sys_link (3)

```
6241 bad:
    ilock(ip);
    ip->nlink++;
    iupdate(ip);
    iunlockput(ip);
    end_op();
    return -1;
}
```

sys_unlink

```
6301 int sys_unlink(void) {  
    struct inode *ip, *dp;  
    struct dirent de;  
    char name[DIRSIZ], *path;  
    uint off;  
  
    if(argstr(0, &path) < 0)  
        return -1;  
    begin_op();  
    if((dp = nameiparent(path, name)) == 0) {  
        end_op();  
        return -1;  
    }  
    ilock(dp);  
    if(namecmp(name, ".")==0 || namecmp(name, "..")==0)  
        goto bad;
```

sys_unlink (2)

6323

```
if((ip = dirlookup(dp, name, &off)) == 0)
    goto bad;
ilock(ip);

if(ip->nlink < 1)
    panic("unlink: _nlink < 1");
if(ip->type == T_DIR && !isdirempty(ip)){
    iunlockput(ip);
    goto bad;
}

memset(&de, 0, sizeof(de));
if(writei(dp, (char*)&de, off, sizeof(de)) != sizeof(de))
    panic("unlink: _writei");
```

sys_unlink (3)

```
6337  if (ip->type == T_DIR) {  
        dp->nlink --;  
        iupdate(dp);  
    }  
    iunlockput(dp);  
  
    ip->nlink --;  
    iupdate(ip);  
    iunlockput(ip);  
  
    end_op();  
  
    return 0;
```

sys_unlink (4)

```
6350 bad :  
    iunlockput(dp);  
    end_op();  
    return    1 ;  
}
```