

The Buffer Layer

xv6-rev7

Carmi Merimovich

CS School, Tel Aviv Academic College

January 12, 2017

Buffer layer routines

- `bp = bread(dev,nblock);`
- `bwrite(bp);`
- `brelease(bp);`
- `nblock = balloc();`
- `bfree(nblock);`

struct buf

```
3501 struct buf {  
    int flags; //B_BUSY, B_VALID, B_DIRTY  
    uint dev;  
    uint sector;  
    struct buf *prev;  
    struct buf *next;  
    struct buf *qnext;  
    uchar data[512];  
};
```

First looks

- We begin by examining:
 - bread.
 - bwrite.
- The driver layer supplies the routine `iderw(bp)`:
 - If `B_VALID` is clear, `iderw` will read the block.
 - If `B_DIRTY` set, `iderw` will write the block.

bread

```
4101 struct buf bread(uint dev, uint sector) {  
    struct buf *b;  
  
    b = bget(dev, sector);  
    if (!(b->flags & B_VALID))  
        bderw(b);  
    return b;  
}
```


bwrite

```
4113 void bwrite(struct buf *b) {  
    if ((b->flags & B_BUSY) == 0)  
        panic(" bwrite");  
    b->flags |= B_DIRTY;  
    iderw(b);  
}
```


The buffer cache

```
4028 struct {  
    struct spinlock lock;  
    struct buf buf[NBUF];  
  
    // Linked list of all buffers, through prev/next.  
    // head.next is most recently used.  
    struct buf head;  
} bcache;
```


Cache list, tail insertion




Cache list, tail insertion




Cache list, tail insertion




Cache list, head insertion





Cache list, head insertion




Cache list, head insertion









```
bp0=bget(1,10);
```




`bp0=bget(1,10);`




```
bp1=bget(1,100);
```




```
bp1=bget(1,100);
```




`breuse(bp1)`




brese(bp1)




```
bp2=bget(1,100);
```




bp2=bget(1,100);




```
bp3=bget(1,20);
```




bp3=bget(1,20);




breuse(bp2)




brealse(bp2);




breuse(bp1)




`brealse(bp1);`




```
bp4=bget(1,200);
```



```
bp4=bget(1,200);
```



```
bp5=bget(1,300);
```



```
bp5=bget(1,300);
```


Cache initialization

⋮

4051

```
// Create linked list of buffers
bcache.head.prev = &bcache.head;
bcache.head.next = &bcache.head;
for (b = bcache.buf; b < bcache.buf+NBUF; b++){
    b->next = bcache.head.next;
    b->prev = &bcache.head;
    b->dev = -1;
    bcache.head.next->prev = b;
    bcache.head.next = b;
}
}
```

bget (1)

```
4065 static struct buf *bget(uint dev, uint sector) {
    struct buf *b;
    acquire(&bcache.lock);
loop:
    // Is the sector already cached?
    for (b = bcache.head.next; b != &bcache.head;
         b = b->next) {
        if (b->dev == dev && b->sector == sector){
            if (!(b->flags & B_BUSY)){
                b->flags |= B_BUSY;
                release(&bcache.lock);
                return b;
            }
            sleep(b, &bcache.lock);
            goto loop;
        }
    }
```

bget (2)

```
// Not cached; recycle some non busy and clean buffers
for (b = bcache.head.prev; b != &bcache.head;
      b = b->prev){
    if ((b->flags & B_BUSY) == 0 &&
        (b->flags & B_DIRTY) == 0){
        b->dev = dev;
        b->sector = sector;
        b->flags = B_BUSY;
        release(&bcache.lock);
        return b;
    }
}
panic("bget: no buffers");
}
```

brelease

```
4124 void brelease(struct buf *b) {  
    if ((b->flags & B_BUSY) == 0)  
        panic("brelease");  
  
    acquire(&bcache.lock);  
  
    b->next->prev = b->prev;  
    b->prev->next = b->next;  
    b->next = bcache.head.next;  
    b->prev = &bcache.head;  
    bcache.head.next->prev = b;  
    bcache.head.next = b;  
  
    b->flags &= ~B_BUSY;  
    wakeup(b);  
    release(&bcache.lock);  
}
```

bitmap

Looking for cleared bit in a byte, the set it.

```
uchar c;
```

```
if ((c & 1) == 0) {c |= 1; return 0};
```

```
if ((c & 2) == 0) {c |= 2; return 1;}
```

```
if ((c & 4) == 0) {c |= 4; return 2;}
```

```
if ((c & 128) == 0) {c |= 128; return 7;}
```

Find cleared bit in a loop

bit 0: 1 << 0

bit 1: 1 << 1

bit 7: 1 << 7

```
for (bi=0; bi < 8; bi++) {  
    m = 1<<bi;  
    if ((c & m) == 0) {  
        c |= m;  
        return (bi);  
    }  
}
```

Large bitmap

```
uchar c[512];

for (bi =0; bi < 4096; bi++) {
    m = 1 << (b i% 8);
    if ((c[bi/8] & m) == 0) {
        c[bi/8] |= m;
        return (bi);
    }
}
```

Even larger bitmap

```
uchar c0[512], c1[512];
```

```
for (bi = 0; bi < 4096; bi++) {  
    m = 1 << (bi % 8);  
    if ((c0[bi/8] & m) == 0) {  
        c0[bi/8] |= m;  
        return (bi);  
    }  
}
```

```
}  
for (bi = 0; bi < 4096; bi++) {  
    m = 1 << (bi % 8);  
    if ((c1[bi/8] & m) == 0) {  
        c1[bi/8] |= m;  
        return (4096+bi);  
    }  
}
```


Really long bitmap

```
uchar c[2][512];

for (b = 0; b < 4096*2; b += 4096) {
    block = b/4096;
    for (bi = 0; bi < 4096; bi++) {
        m = bi % 8;
        if ((c[block][bi/8] & m) == 0) {
            c[block][bi/8] |= m;
            return (b+bi);
        }
    }
}
```

Bitmap not necessarily multiple of 4096

```
#define NUMOFBITS 8000
```

```
uchar c[2][512];
```

```
for (b = 0; b < NUMOFBITS; b += 4096) {  
    block = b/4096;  
    for (bi = 0; bi < 4096 && b+bi < NUMOFBITS; bi++)  
        if ((c[block][bi/8] & m) == 0) {  
            c[block][bi/8] |= m;  
            return (b+bi);  
        }  
}
```

balloc

```
4453 static uint balloc(uint dev) {  
    struct superblock sb;  
    struct buf *bp = 0;  
    readsb(dev, &sb);  
    for (int b = 0; b < sb.size; b += BPB){  
        bp = bread(dev, BBLOCK(b, sb.ninodes));  
        for (int bi = 0; bi < BPB && b + bi < sb.size; bi++) {  
            int m = 1 << (bi % 8);  
            if ((bp->data[bi/8] & m) == 0){ // Is block free?  
                bp->data[bi/8] |= m; // Mark block in use.  
                log_write(bp);  
                brelse(bp);  
                bzero(dev, b + bi);  
                return b + bi;  
            }  
        }  
        brelse(bp);  
    }  
    panic("balloon out of blocks");  
}
```

bitmap, clearing bit

```
uchar c;
```

```
m = 1 << bi;
```

```
c &= ~m;
```

bitmap, clearing bit

```
uchar c[512];
```

```
m = 1 << (bi % 8);
```

```
c[bi/8] &= ~m;
```

bitmap, clearing bit

```
uchar c[2][512];
```

```
block = b/4096;
```

```
bi = b % 4096;
```

```
m = 1 << (bi % 8);
```

```
c[block][bi/8] &= ~m;
```

bfree

```
4480 static void bfree(int dev, uint b) {  
    struct buf *bp;  
    struct superblock sb;  
    int bi, m;  
  
    readsb(dev, &sb);  
    bp = bread(dev, BBLOCK(b, sb.ninodes));  
    bi = b % BPB;  
    m = 1 << (bi % 8);  
    if ((bp->data[bi/8] & m) == 0)  
        panic("freeing _free_block");  
    bp->data[bi/8] &= ~m;  
    log_write(bp);  
    brelse(bp);  
}
```